




# Introduction to Storage Options

Block, File and Object Storage

# Block Storage



Block storage divides data into fixed-sized blocks and stores them as separate pieces

- Data divided into fixed-size blocks
- Unique identifier
- Retrieve data by its block ID

**Low Latency:** high-performance applications requiring low latency.

**High Customizability:** fine-tuned storage optimizations (e.g., databases, transactional applications).

**Compatibility:** high-performance applications like virtual machines, databases, and enterprise applications.

**Complex Management:** Block storage can be harder to manage, especially when scaling.

**No Metadata:** Blocks don't store metadata, so context for each block isn't preserved.


# File Storage



Amazon Elastic File System  
(Amazon EFS)



Amazon FSx for  
Windows File Server



File storage organizes data in a hierarchical structure of files and folders. Data is stored and retrieved as complete files using a path-based system.

- Hierarchical file structure
- Can define tags and permissions on files
- Enables file sharing and collaboration

**Familiar Structure:** Simple, easy-to-use folder-based structure

**Supports Metadata:** Allows tagging and permissions on files, making data easier to manage and control.

**Shared Access:** Works well in environments where multiple users need to access files (e.g., NAS)

**Scalability Limits:** Less scalable than object storage due to limitations in the hierarchical structure

**Lower Performance for Large Data Sets:** Can be less efficient than block storage for high-performance needs.

# Object Storage



Amazon Simple Storage  
Service (Amazon S3)



Object Storage

File Flat Namespace

Object storage manages data as individual objects with unique identifiers and metadata, storing them in a flat namespace in a repository. Each object includes data, metadata, and a unique ID.

- Data is stored with metadata and a unique ID
- Stored across a distributed network offering durability
- Erasure Coding
- Replication and horizontal scaling support

**Scalability:** Highly scalable, ideal for massive amounts of unstructured data

**Metadata-Rich:** Extensive metadata support, useful for organizing, indexing, and searching data

**Durability and Redundancy:** Often includes built-in replication and redundancy, improving data resilience.

**Latency:** Higher latency than block storage, making it less suited for high-performance applications.

**Limited Compatibility:** Not ideal for applications that require frequent updates to data, such as databases.



# Introduction to Amazon S3

Buckets and Objects

# What is Amazon S3

- **Global Object Storage** Solution – *Accessible globally but hosted regionally.*
- Data is stored as **objects** (with metadata and a globally unique identifier).
- Public service with multi-connect access.
- Use cases include **data lakes**, **websites**, **mobile apps**, **backups**, archives, and **big data analytics**.
- Designed to offer **99.999999999 (11 9s)** of durability.
- From a single CSV file to a high-resolution video.

Object Storage – Not File or Block Storage



Amazon Simple Storage Service (Amazon S3)



Buckets are private by default



IAM Policies



Bucket Policies



ACLs

Access Methods




# What is Amazon S3 – Buckets and Objects


AWS never replicates your data to other regions unless you configure it to do so.


Increases Durability & Availability

Compliance & Data Sovereignty



- Globally unique names across AWS accounts
- 3 to 63 characters in length
- Lowercase letters, numbers, dots (.) and hyphens (-)
- Cannot start with an IP address
- 100 soft limit and 1000 hard limit per account



- Objects
  - Size from 0 bytes to 5TB
  - Key: blueberry-muffin.jpg
  - Value: 

## Flat File Structure

- blueberry-muffin-recipe.doc
- chocolate-muffin-recipe.doc
- lemon-drizzle-muffin-recipe.doc
- /archive/walnut-cake-v2-recipe.doc
- /latest/walnut-cake-v3-recipe.doc

Delimiter (/)

Prefix

The prefix and delimiter parameters can be used to limit the results of a list operation

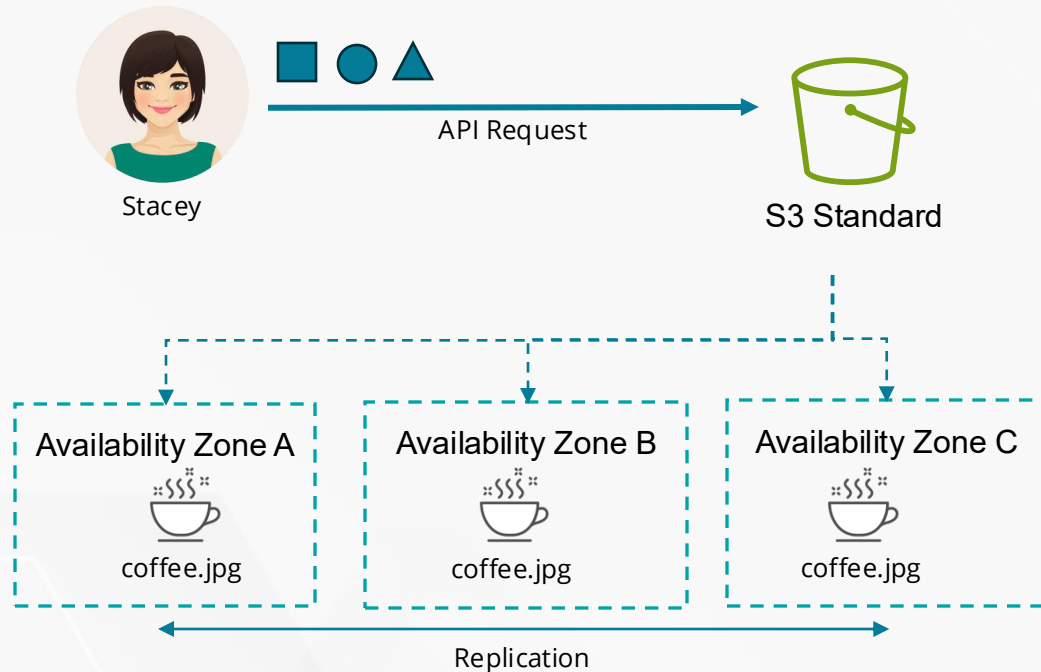


# Amazon S3 Storage Classes

Choosing the best storage class for your  
use case



# Storage Class – S3 Standard




Default Storage Class

Instant Access, milliseconds first byte latency

Idea for frequently accessed data that is not replaceable

Buckets can be made publicly available




Using permissions and/or static website hosting

## Durability

- Protection against data loss
- Prevent data corruption
- 99.999999999% (11 9s) durability

Content-MD5 Checksums and Cyclic Redundancy Checks (CRC) for data integrity



Tom

if you store 10,000,000 objects in Amazon S3, then on average, you can expect to experience the loss of a single object every 10,00 years

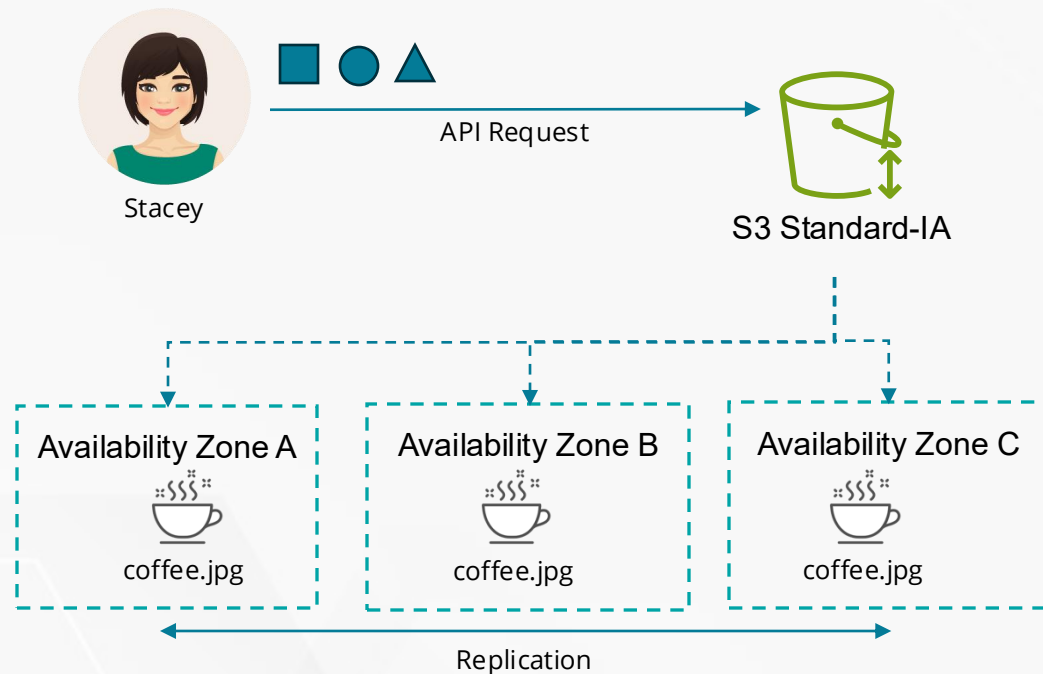


Lisa - CFO

### Costs based on:

- Storage fee GB/month for data stored
- Transfer OUT fee
- Request fee per 1,000 requests

# Storage Class – S3 Standard-IA (Infrequent Access)



Cheaper than S3 Standard

Instant Access: milliseconds first byte latency

Idea for infrequently accessed data that is not replaceable

## Durability

- Protection against data loss
- Prevent data corruption
- 99.999999999% (11 9s) durability
- 99.9% availability

Content-MD5 Checksums and Cyclic Redundancy Checks (CRC) for data integrity

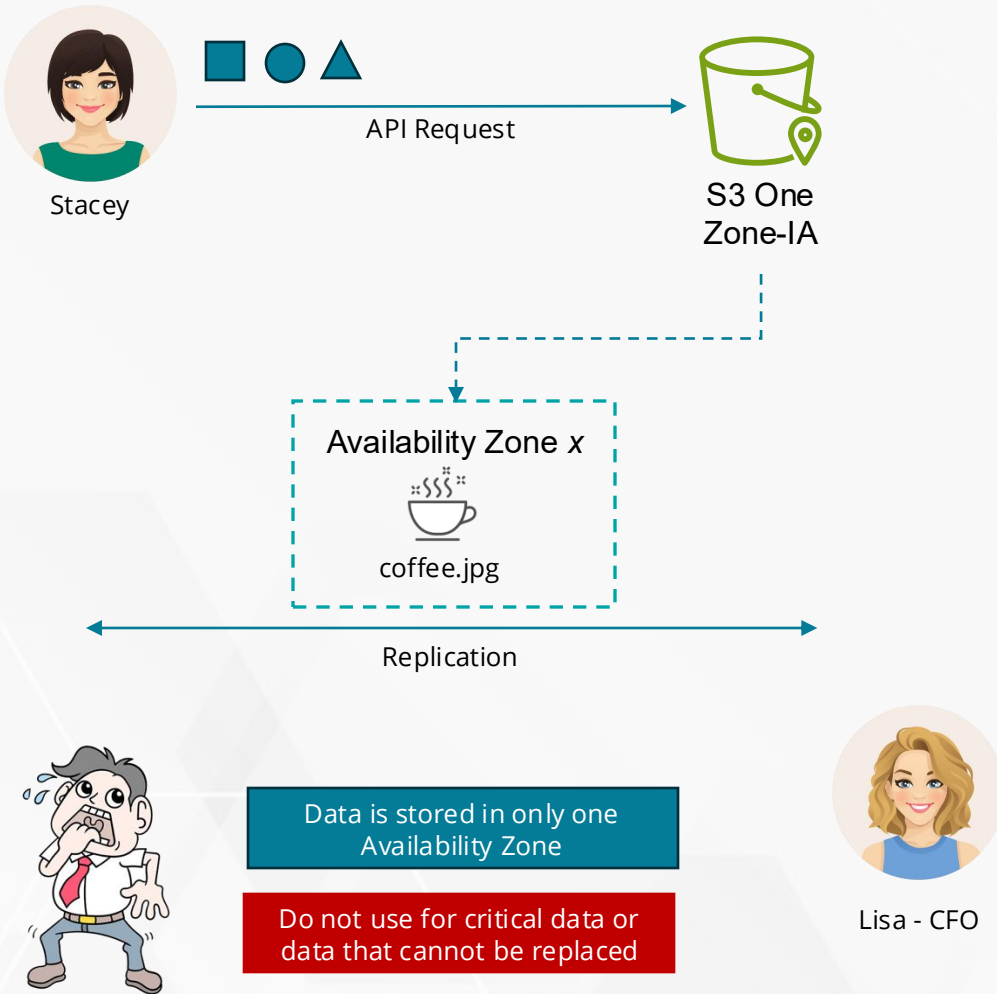
## Costs based on:

- Storage fee GB/month for data stored
- Transfer OUT fee
- Request fee per 1,000 requests
- **Minimum storage charge of 128KB**
- **Minimum duration charge of 30 days**
- **Per GB retrieval cost**



Lisa - CFO

# Storage Class – S3 One-Zone IA



Cheaper than S3 Standard and Standard IA

Instant Access: milliseconds first byte latency

Ideal for infrequently accessed data that is replaceable, e.g., secondary copies or replicas

## Durability

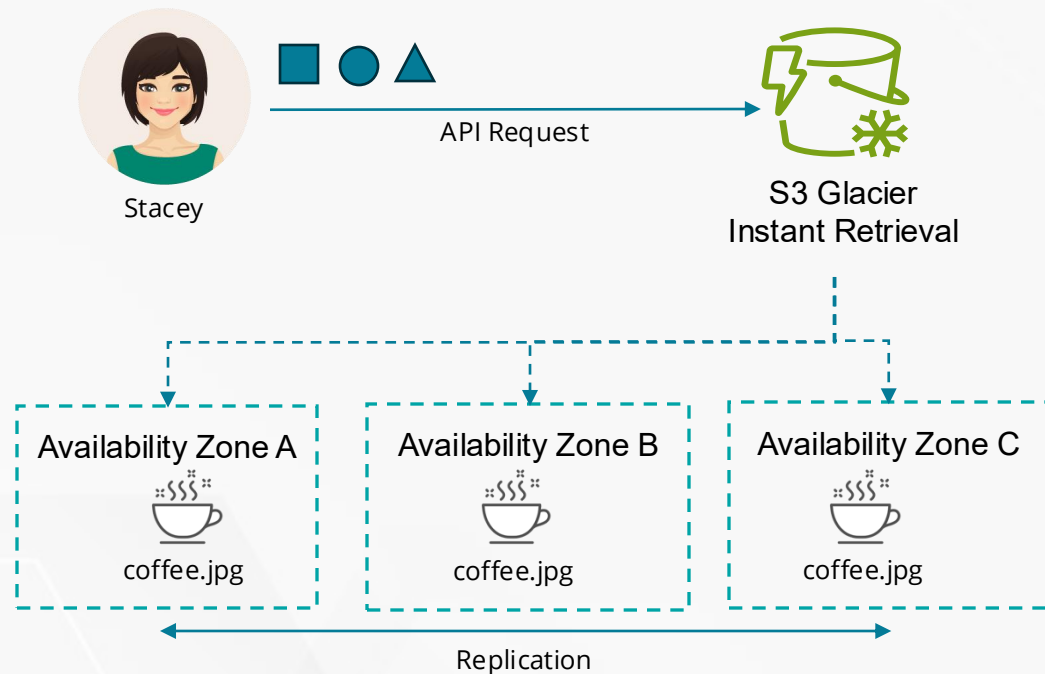
- Protection against data loss
- Prevent data corruption
- 99.999999999% (11 9s) durability
- 99.5% availability

Content-MD5 Checksums and Cyclic Redundancy Checks (CRC) for data integrity

## Costs based on:

- Storage fee GB/month for data stored
- Transfer OUT fee
- Request fee per 1,000 requests
- **Minimum storage charge of 128KB**
- **Minimum duration charge of 30 days**
- **Per GB retrieval cost**

# Storage Class – S3 Glacier – Instant Retrieval



Designed for archival data  
Cheaper than Standard Classes

Instant Access: milliseconds  
first byte latency

Ideal for infrequently accessed data that is not replaceable, e.g. once per quarter

## Durability

- Protection against data loss
- Prevent data corruption
- 99.999999999% (11 9s) durability
- 99.9% availability

Content-MD5 Checksums and  
Cyclic Redundancy Checks  
(CRC) for data integrity

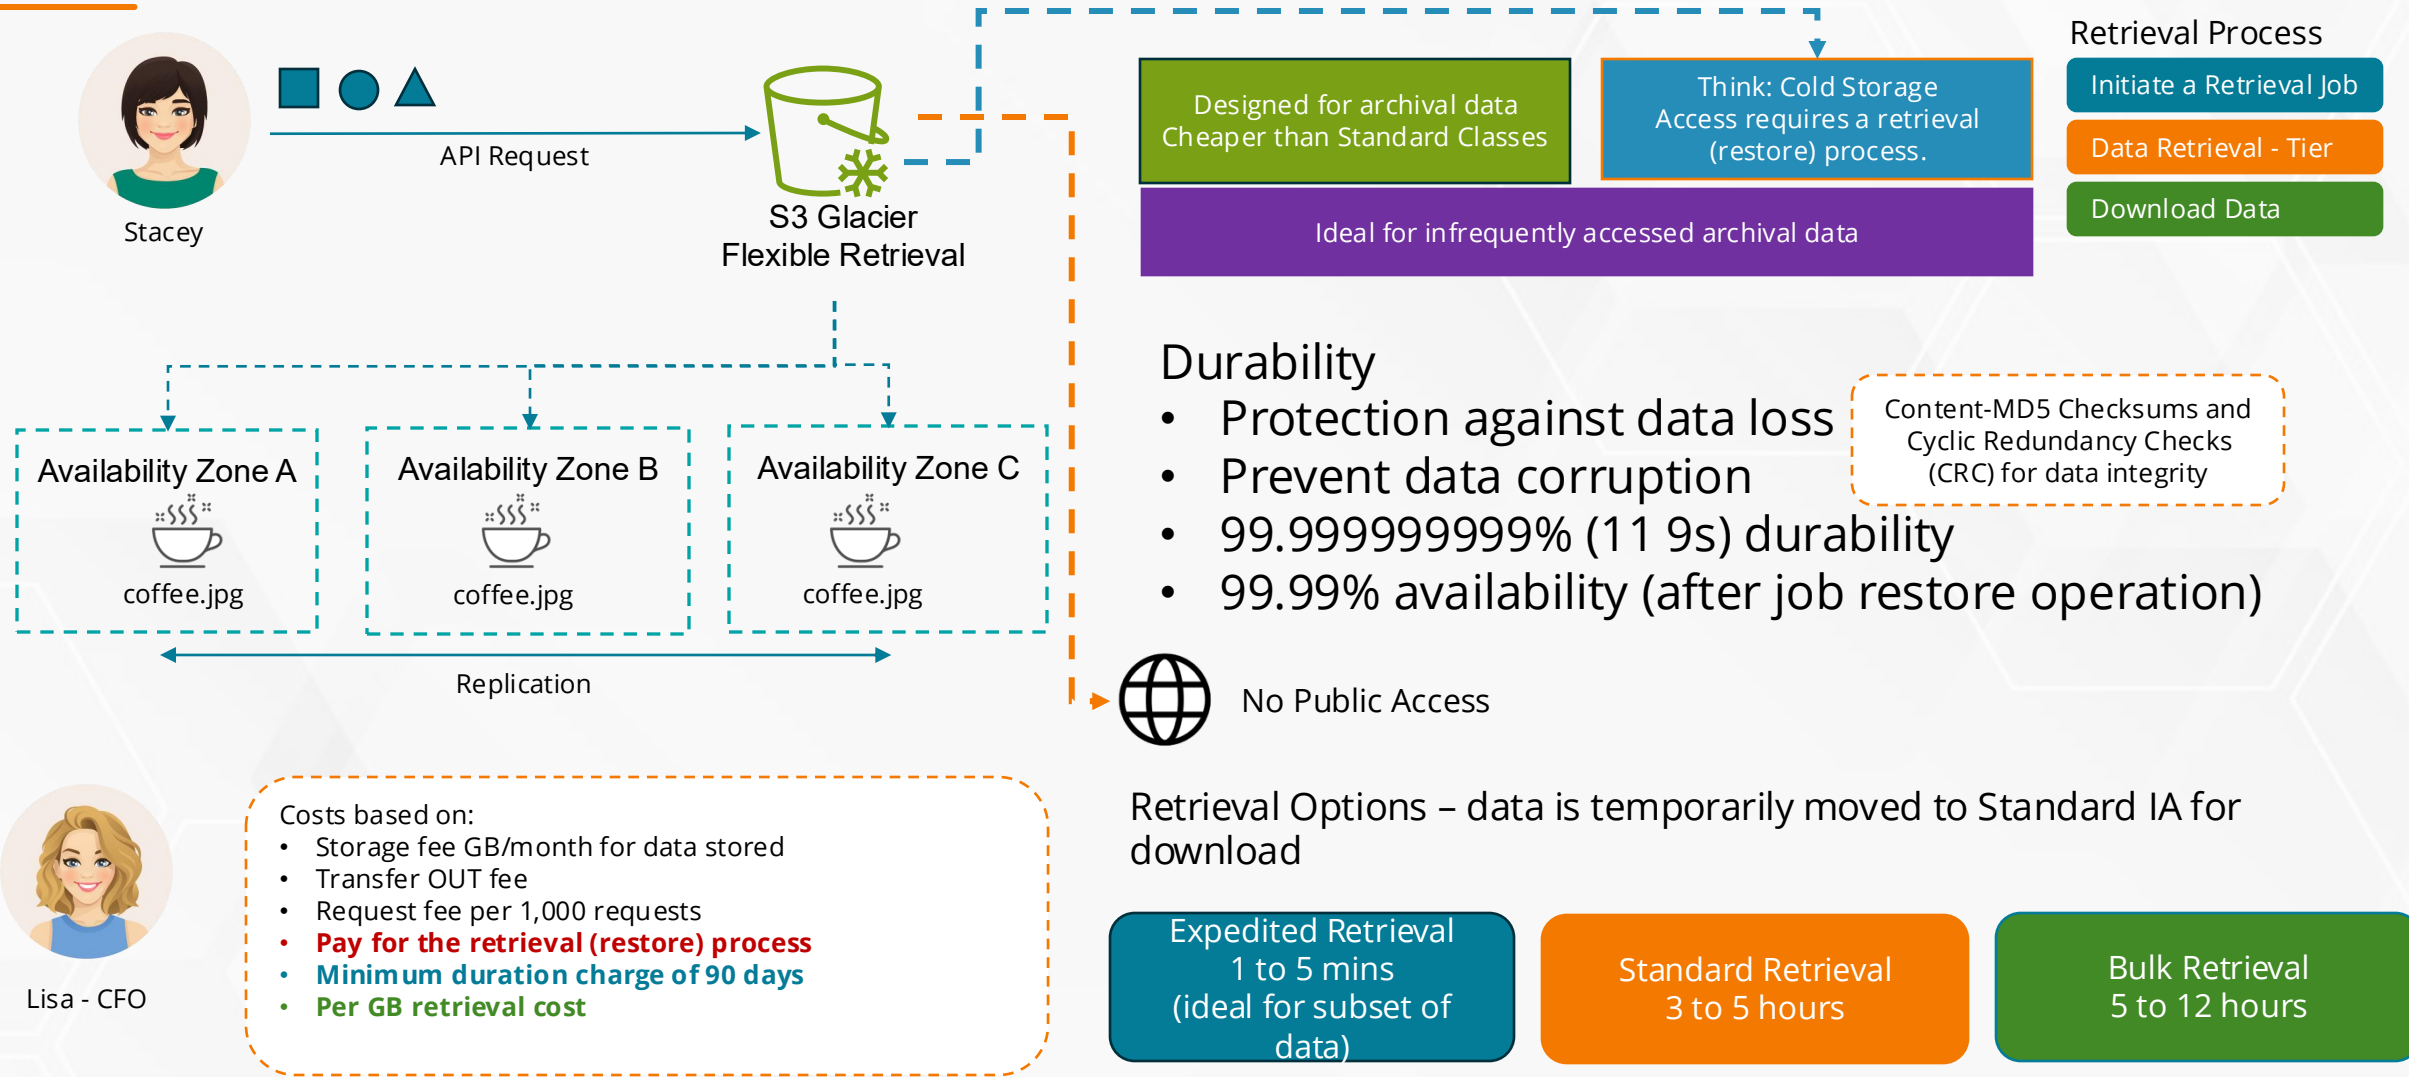
## Costs based on:

- Storage fee GB/month for data stored
- Transfer OUT fee
- Request fee per 1,000 requests
- **Minimum storage charge of 128KB**
- **Minimum duration charge of 90 days**
- **Per GB retrieval cost**

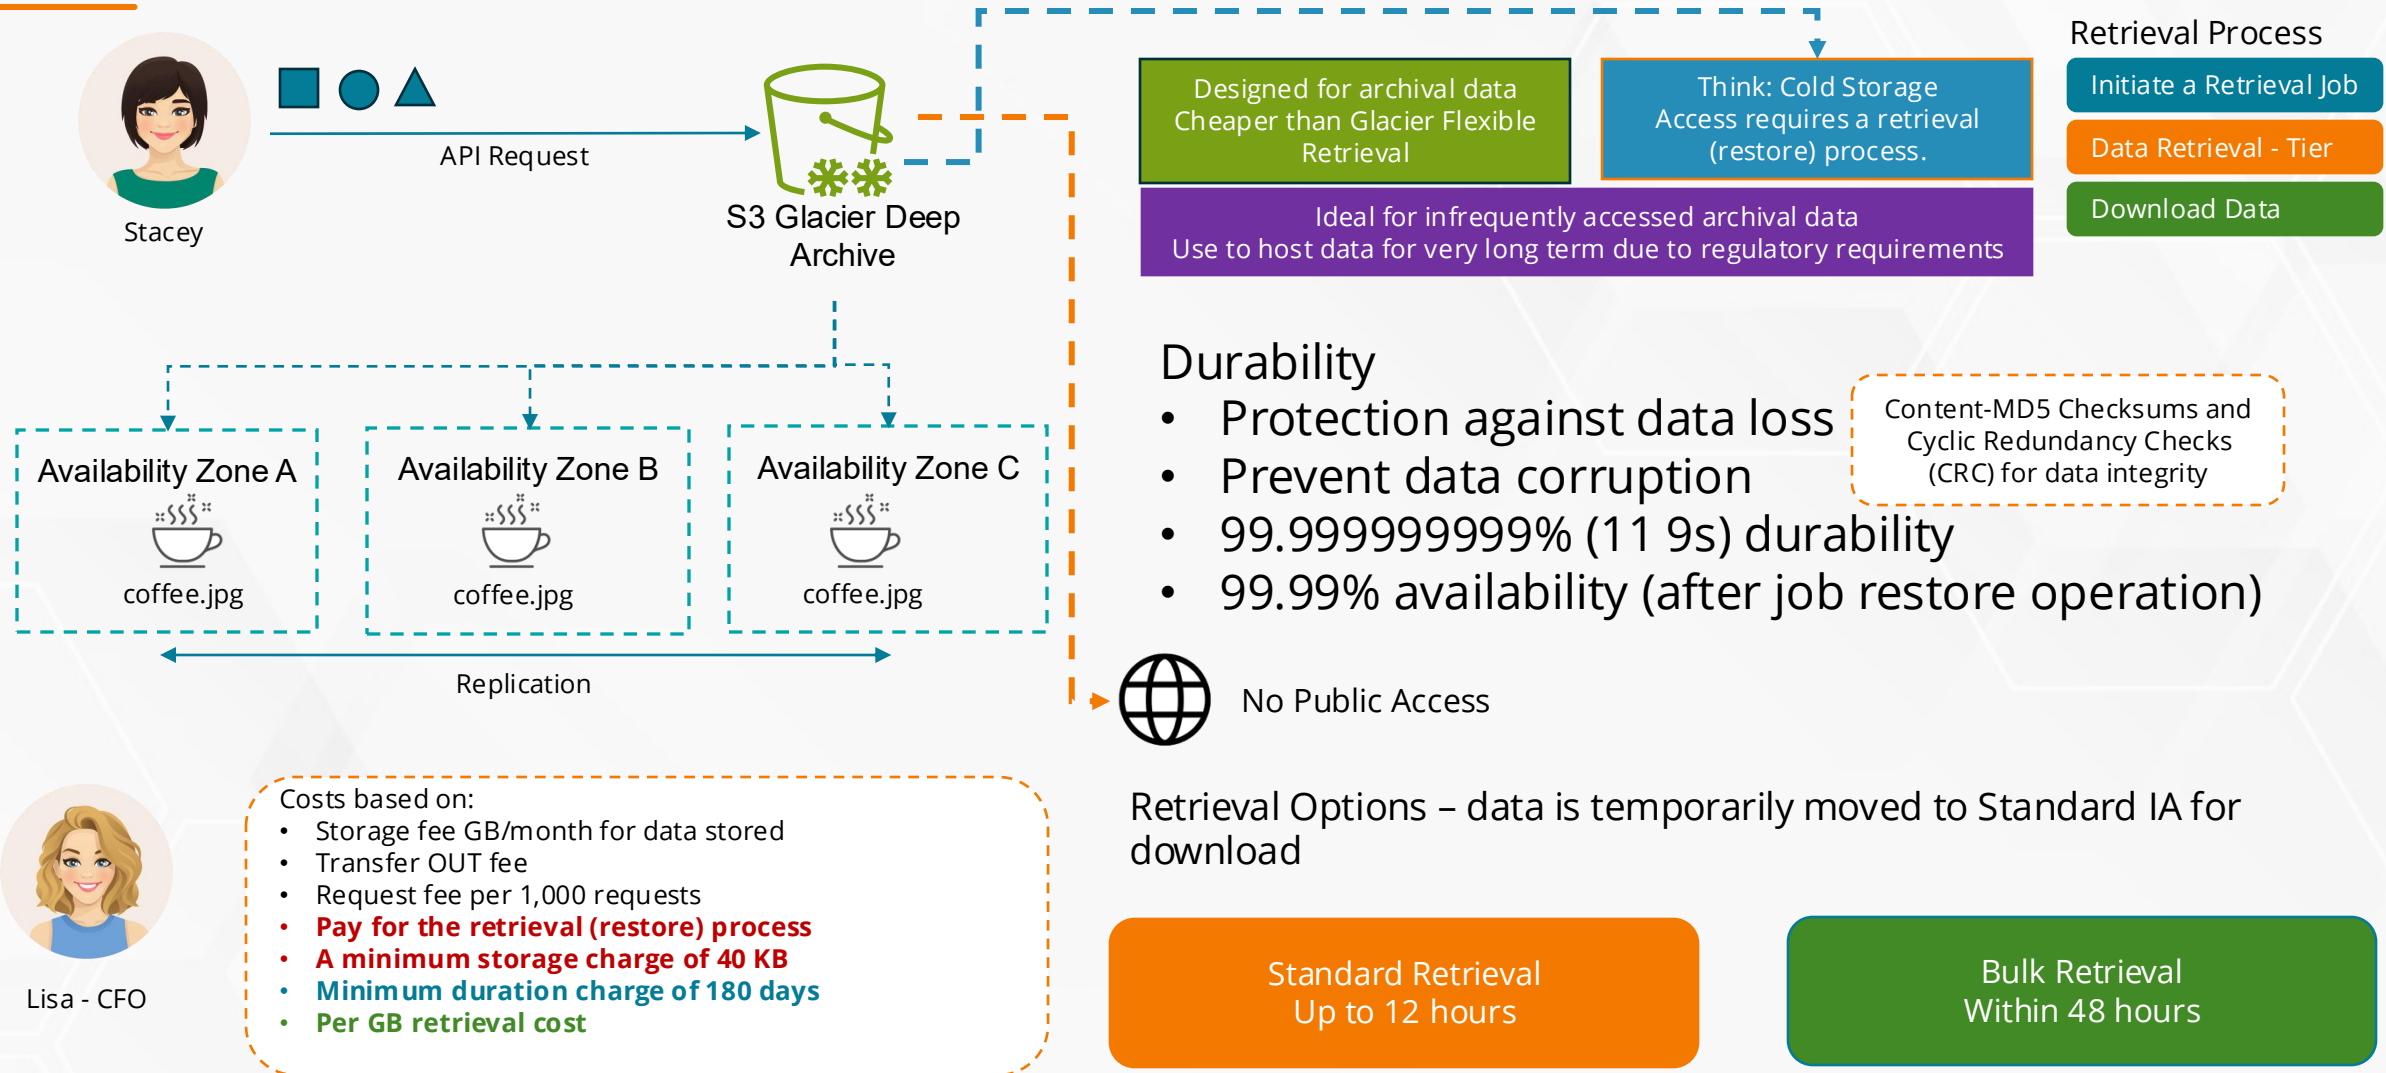


Lisa - CFO

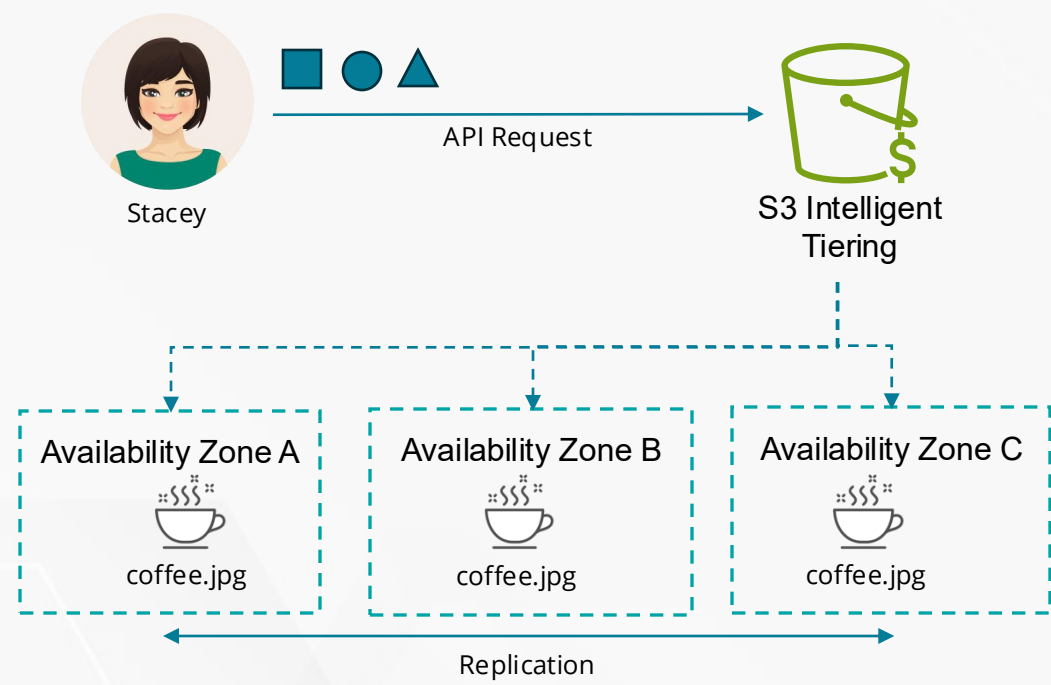
# Storage Class – S3 Glacier – Flexible Retrieval



# Storage Class – S3 Glacier – Deep Archive



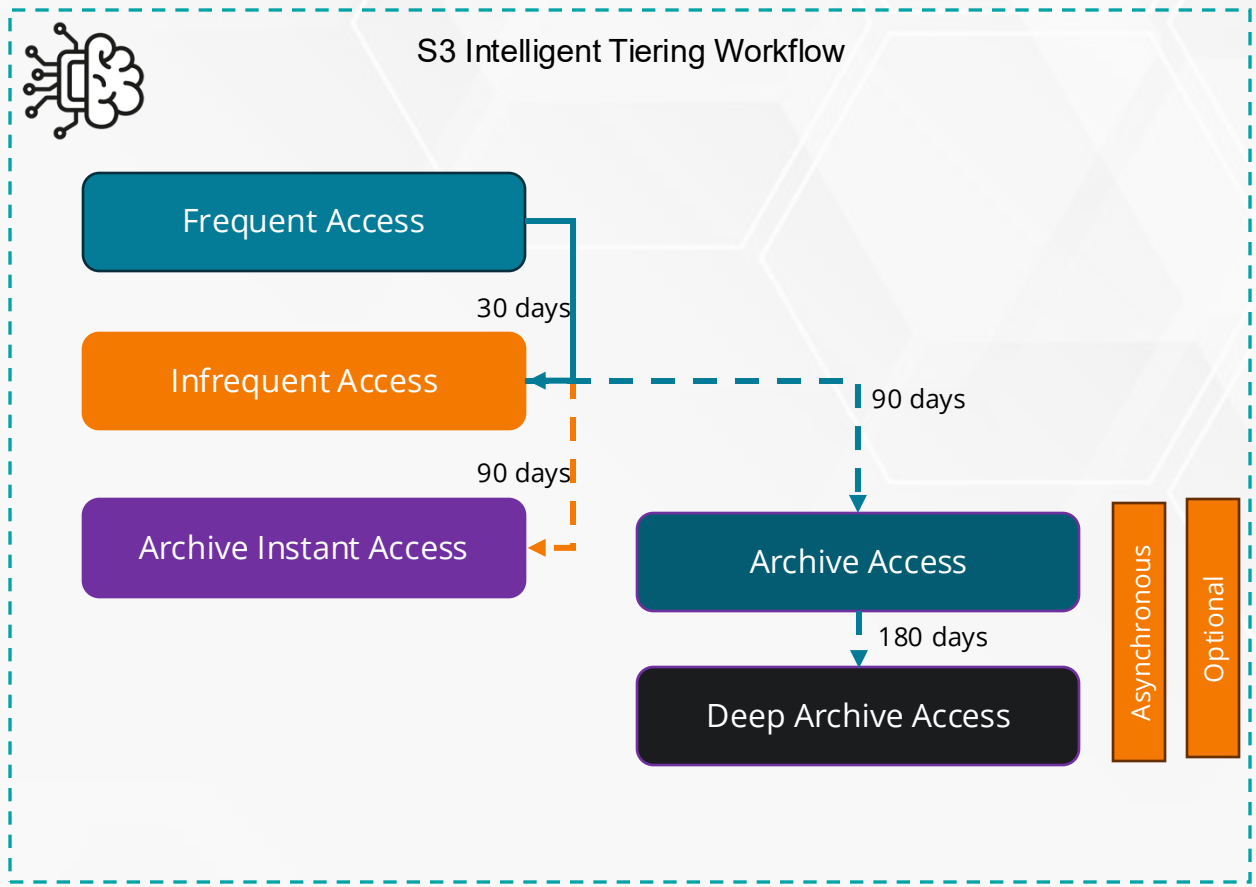
# Storage Class – S3 Intelligent Tiering



Lisa - CFO

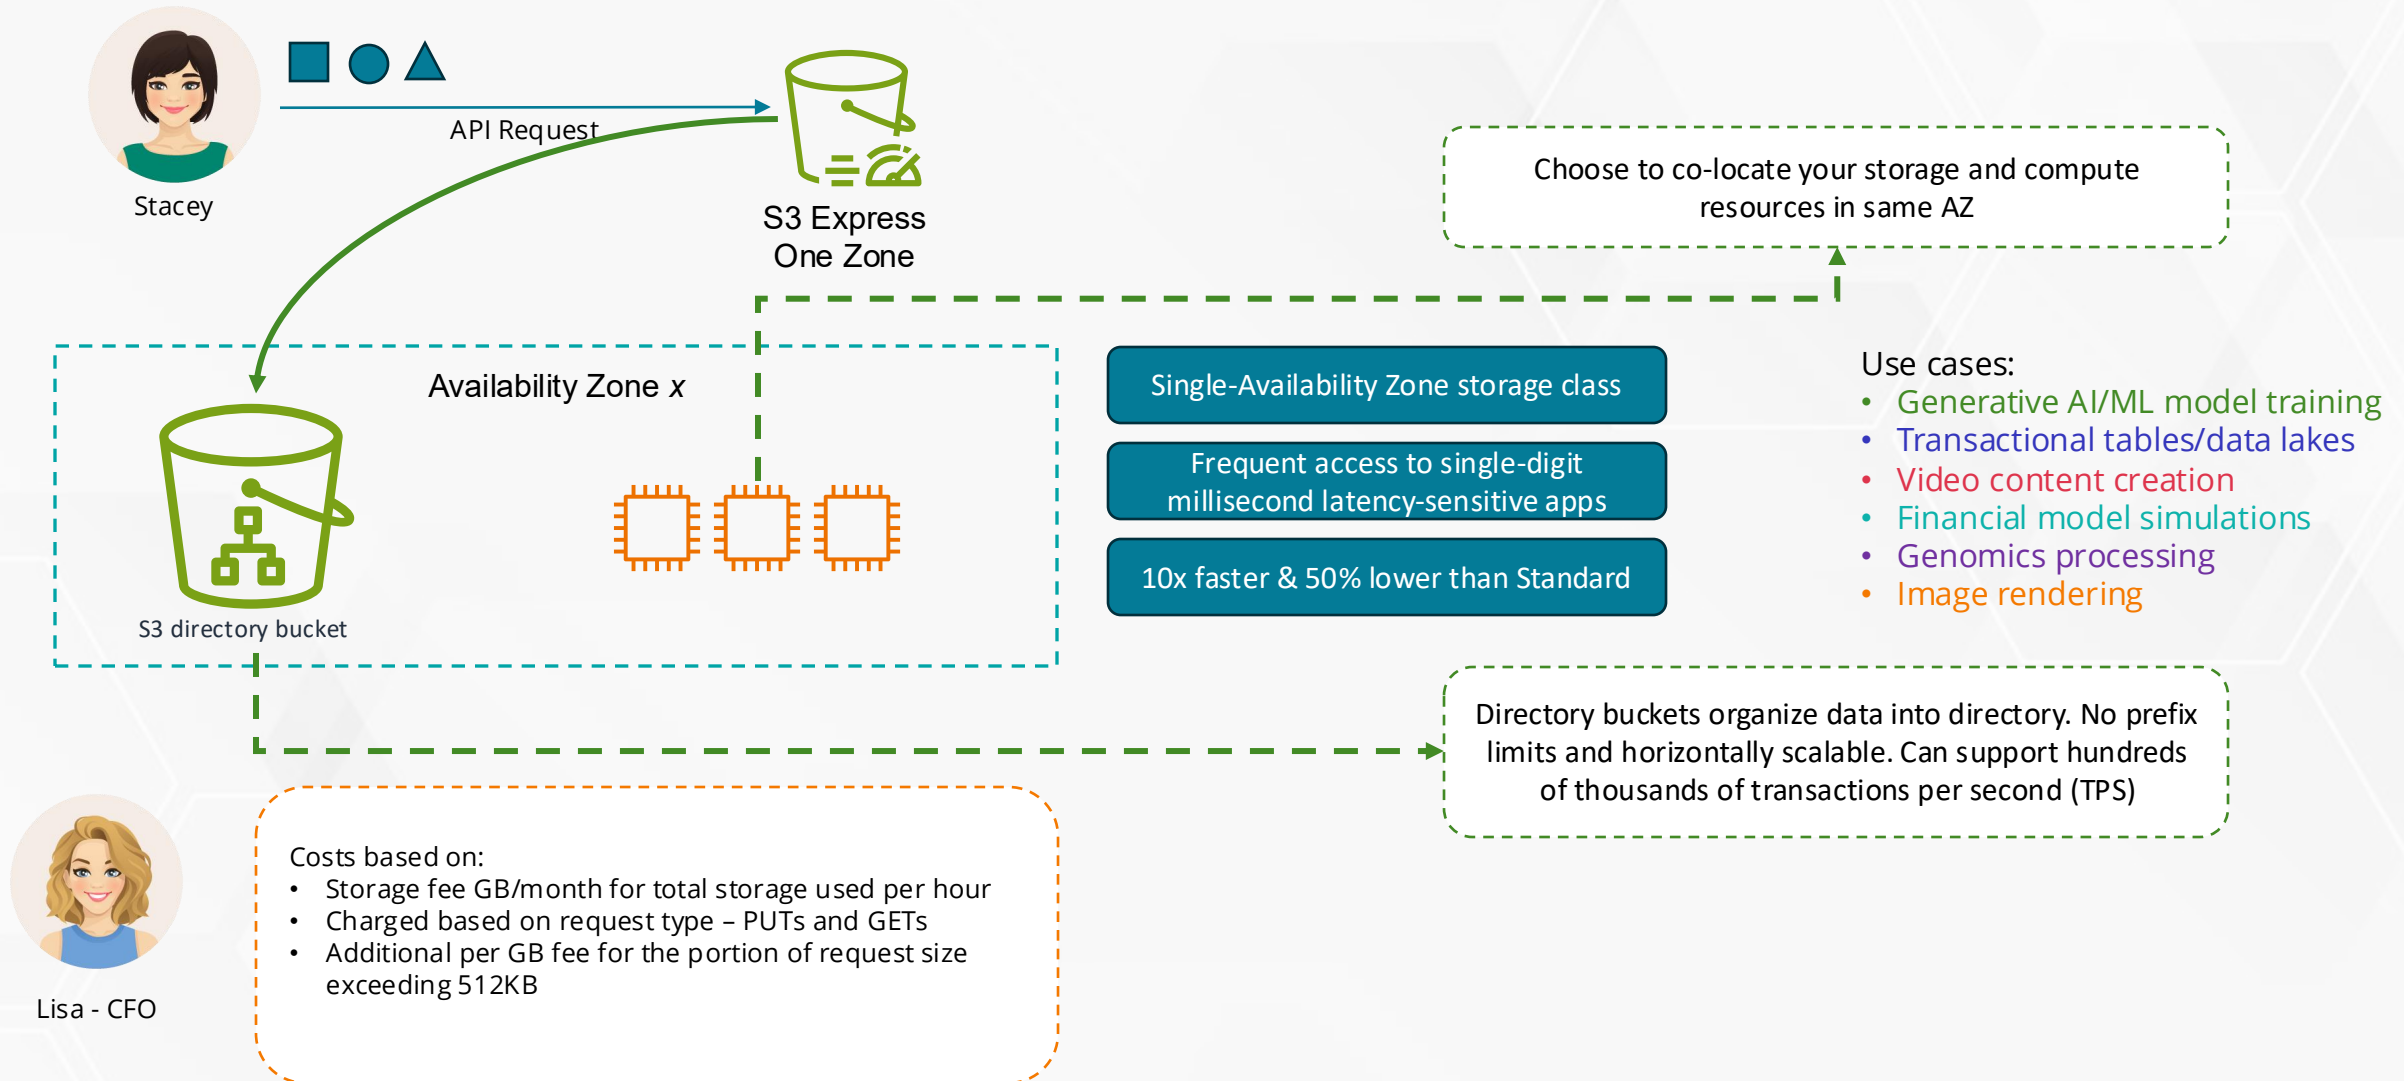
- Costs based on:
- Storage fee GB/month for data stored
  - Transfer OUT fee
  - Request fee per 1,000 requests
  - **Monitoring and automation fee per object**
  - **No retrieval costs**

Designed for data where the access pattern is unknown or can change quickly. Objects not accessed for 30 days are moved into low cost infrequent tier and further into archive instant access, archive access and deep archive access. If the object is accessed again it is moved into Standard tier.





# Storage Class – S3 Express One-Zone







# Create an Amazon S3 Bucket


Features and Configuration



# Amazon S3 Object Versioning

Prevent accidental deletions and overwrites

# Object Versioning



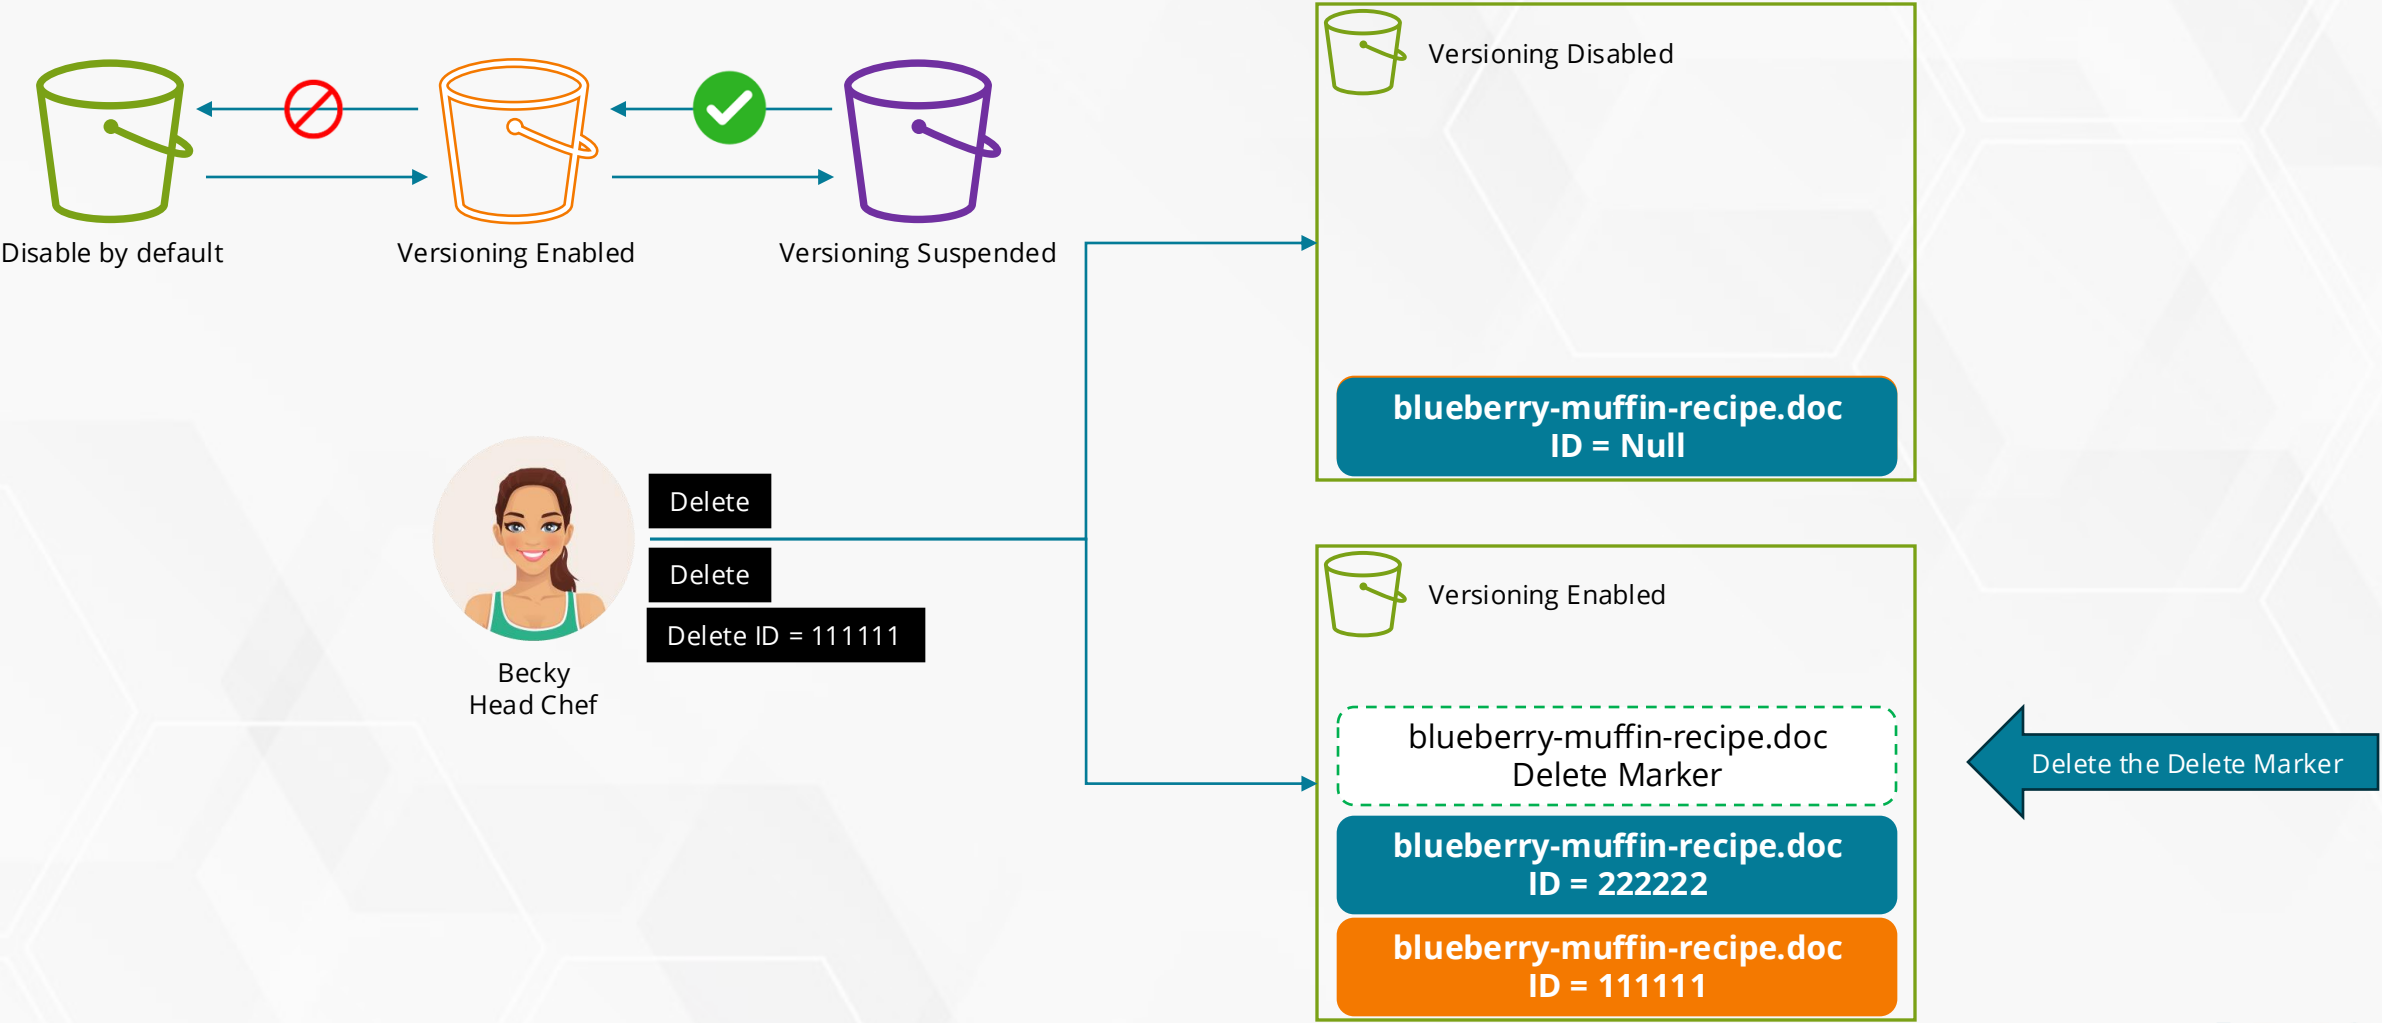
blueberry-muffin-recipe.doc



blueberry-muffin-recipe.doc

Amazon S3 Buckets  
Default: Versioning Disabled

# Bucket Versioning States





# Amazon S3 MFA Delete

Using MFA Delete to prevent against  
accidental deletions

# Bucket Versioning States – Enabling MFA Delete



Becky  
Head Chef

Delete ID = 111111



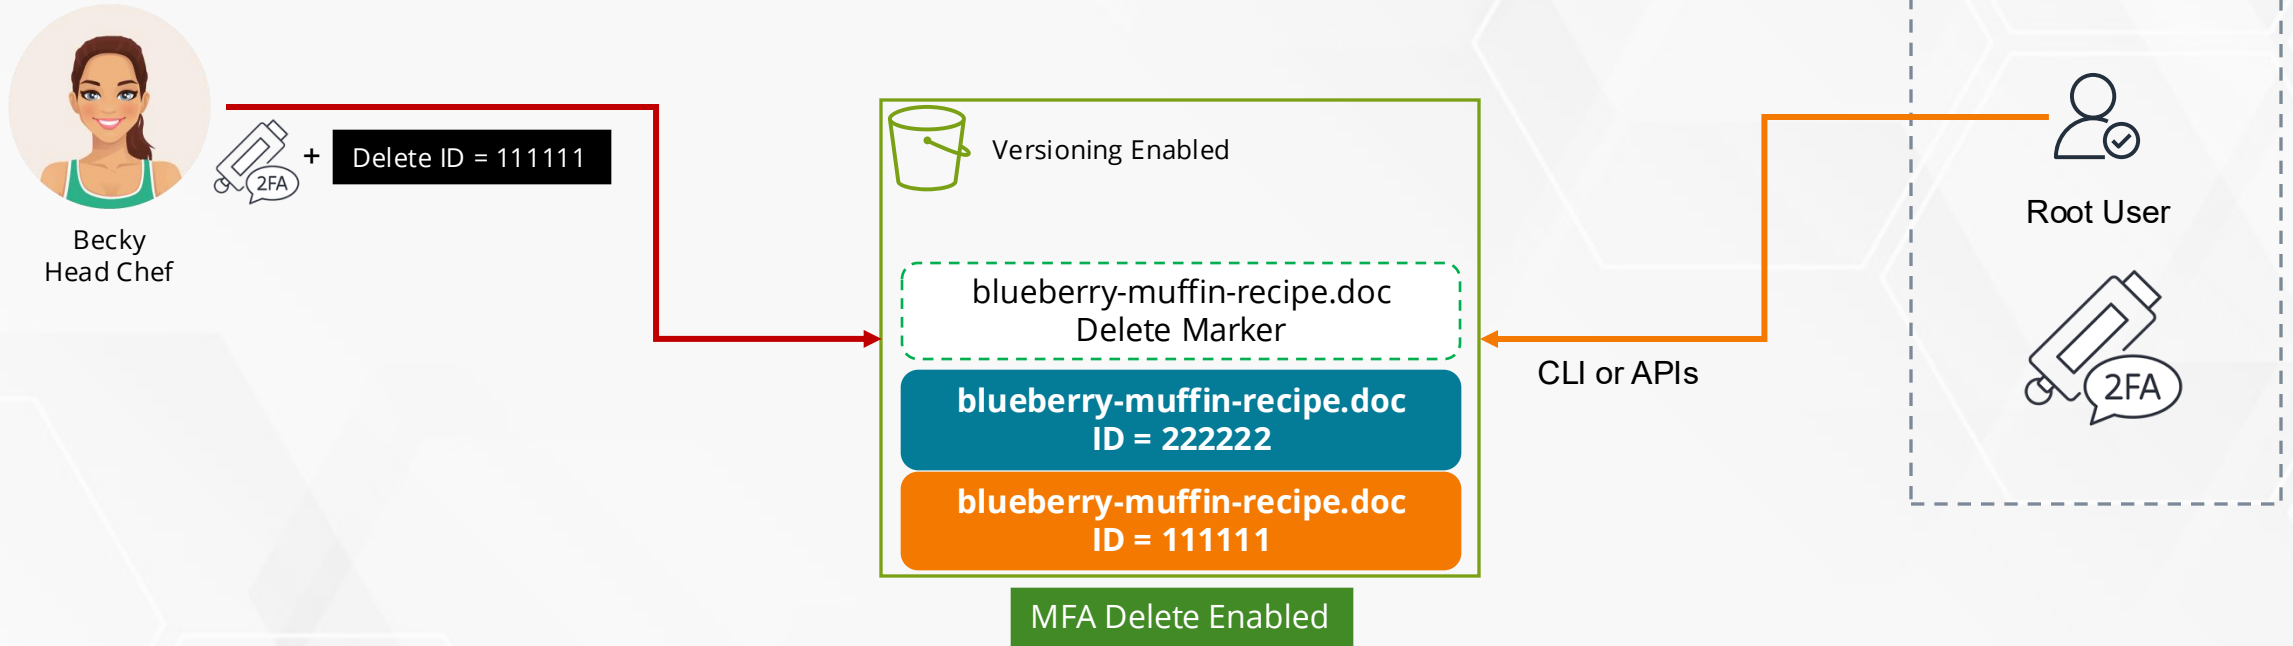
MFA Delete Enabled



Root User



# Bucket Versioning States – Enabling MFA Delete



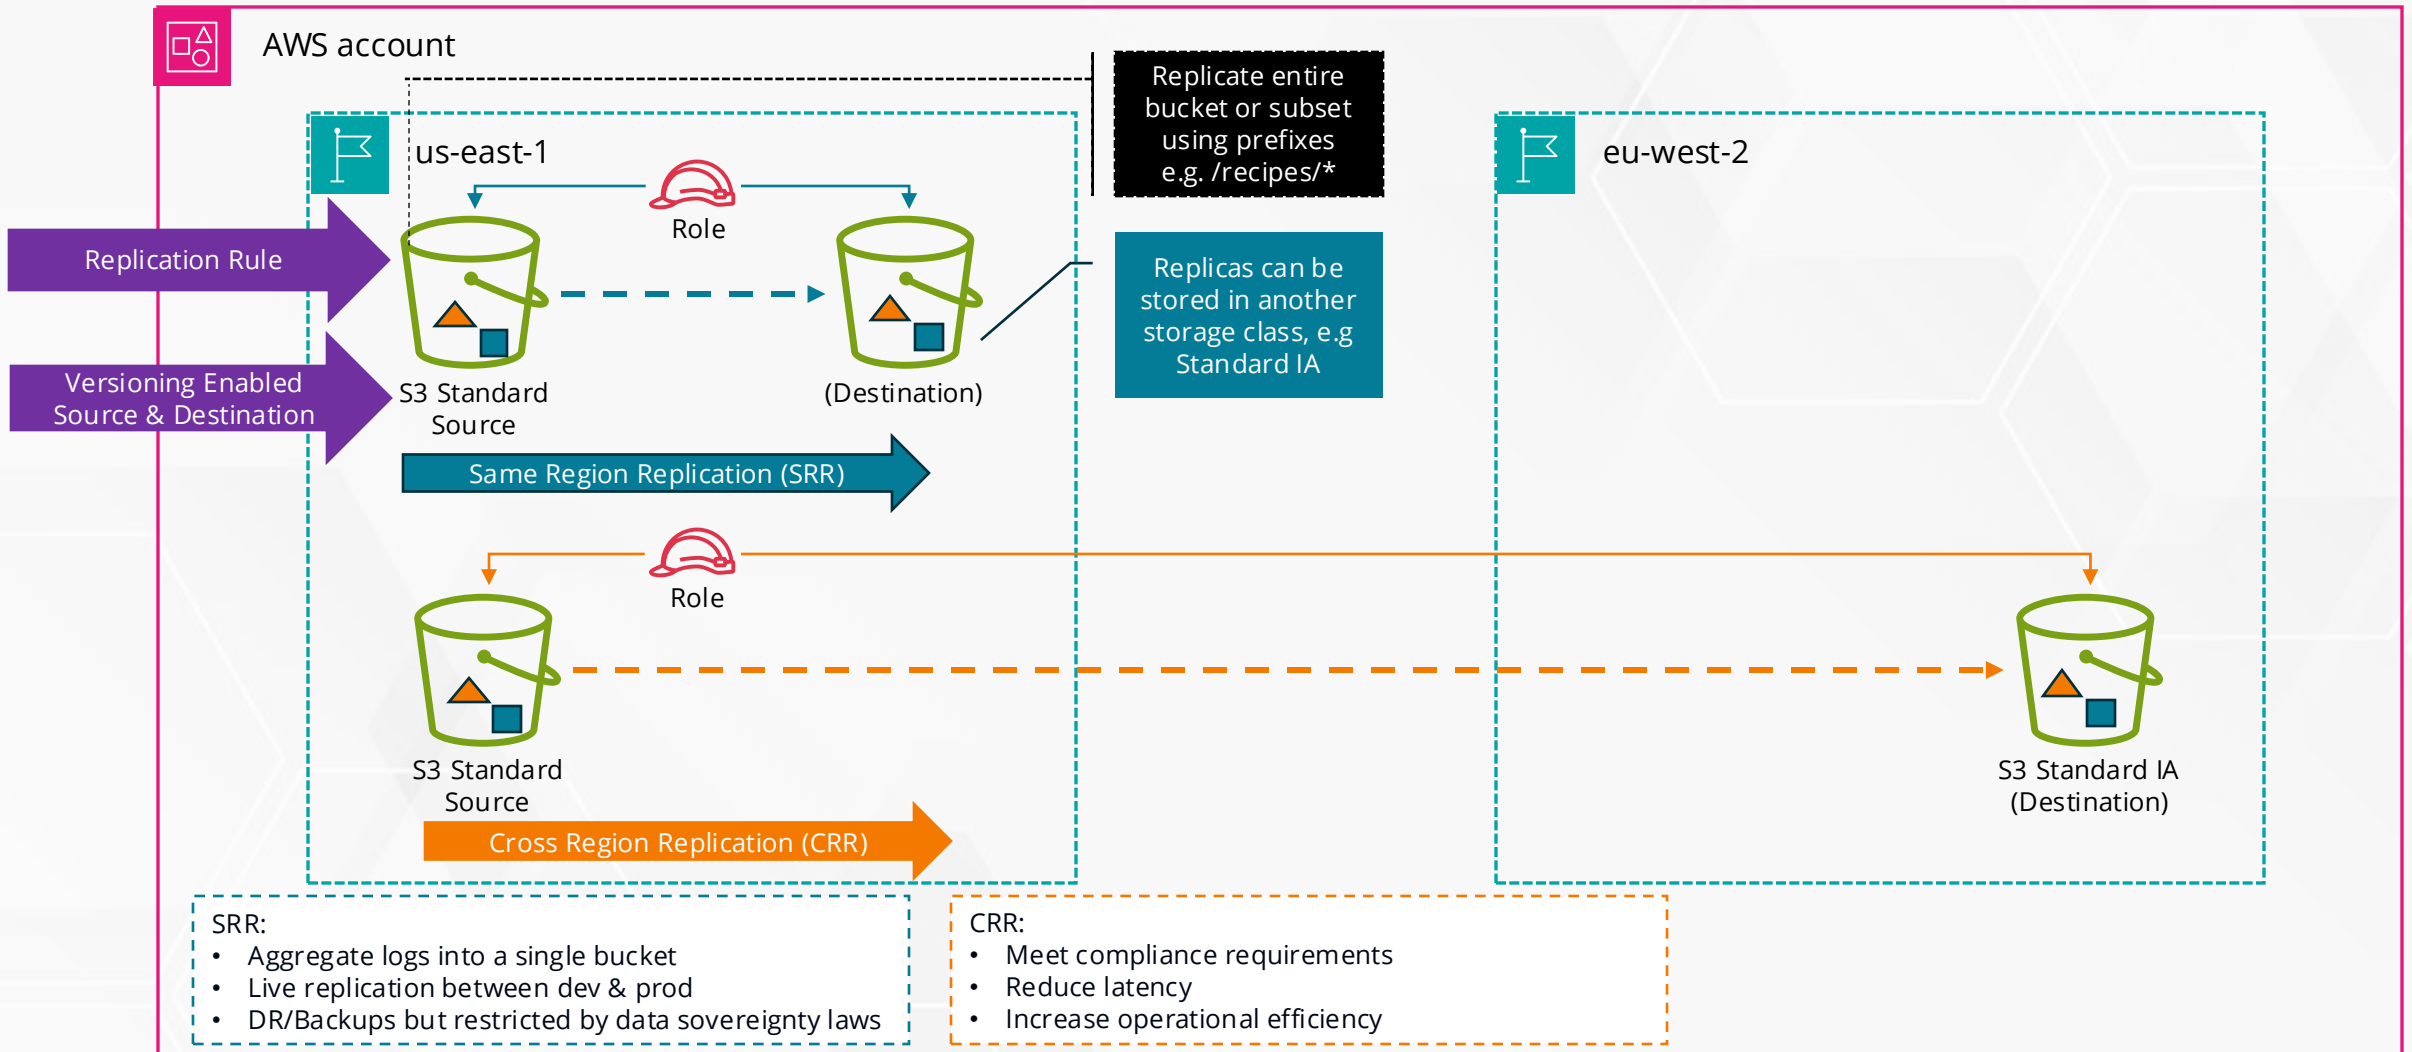


# Amazon S3 Bucket Replication

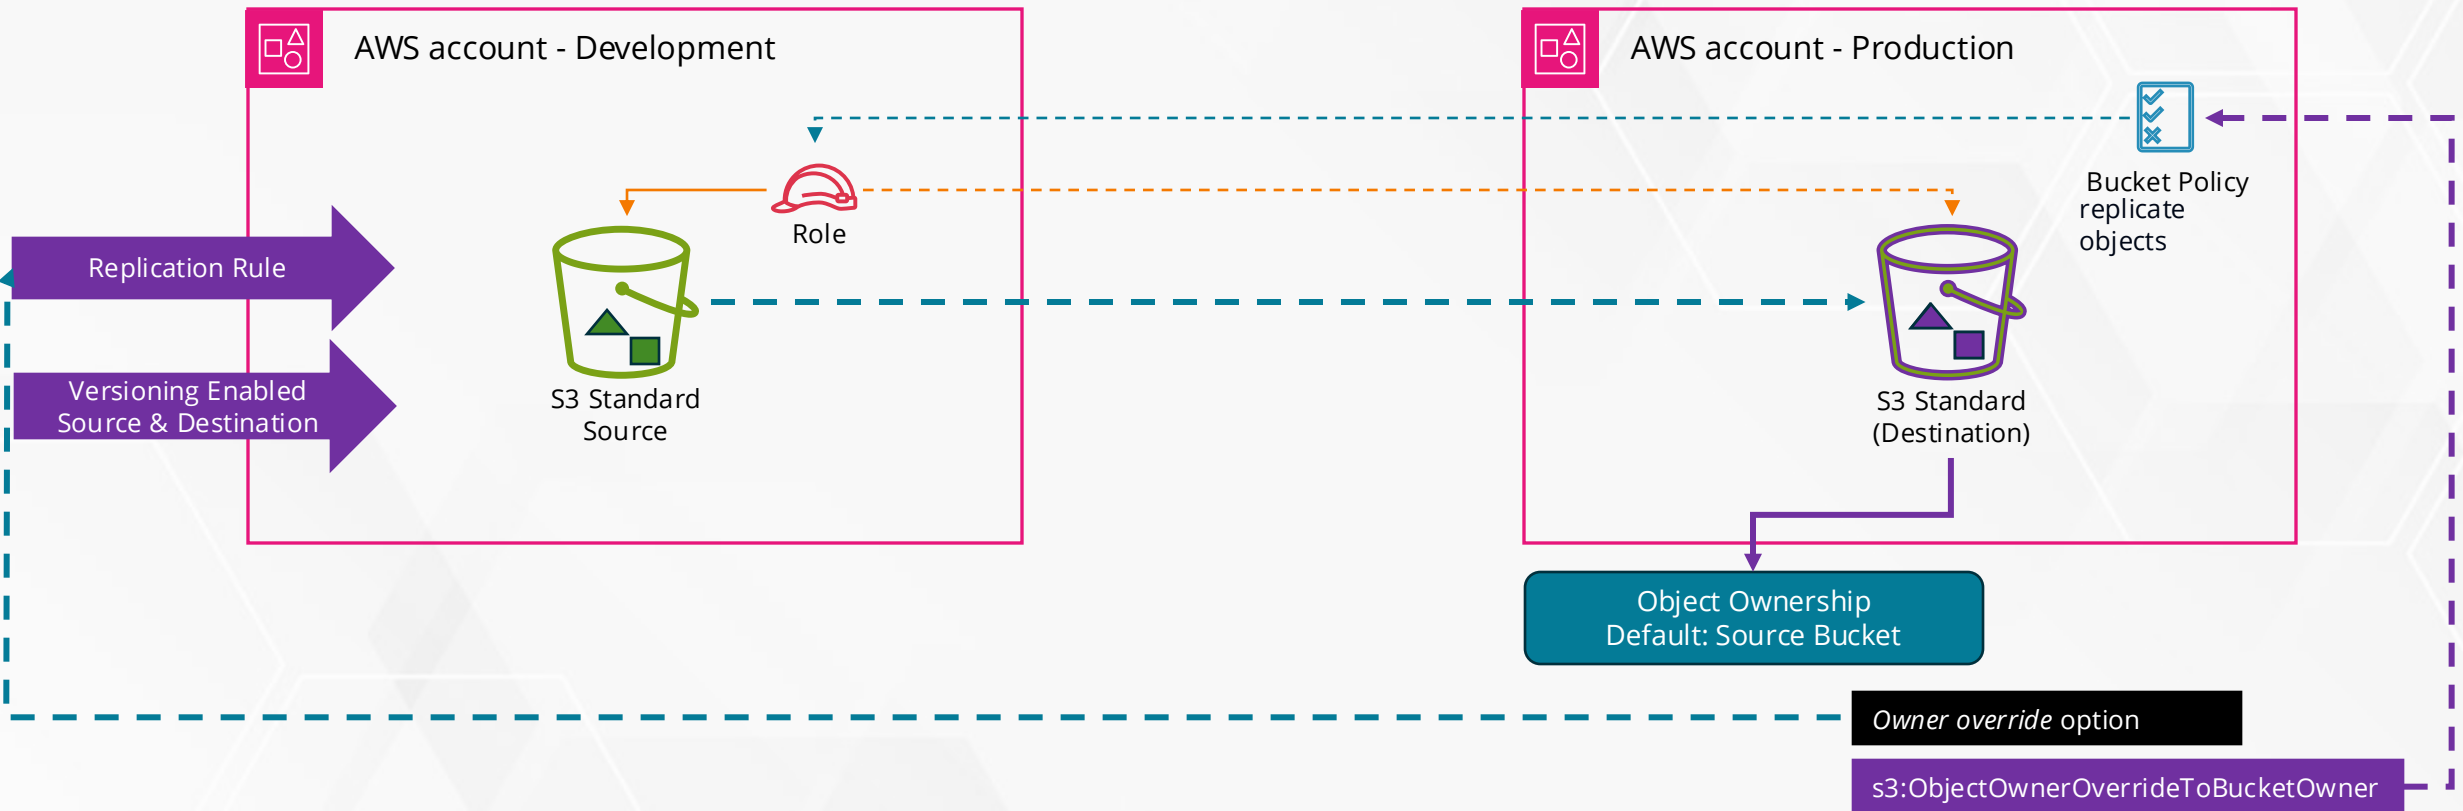
Cross-Region and Same Region Replication




# Amazon S3 Bucket Replication



# Amazon S3 Bucket Replication – Cross Account



# Amazon S3 Replication – Additional Points

- Replication Time Control (**RTC**) – **15 Minutes** Guarantee Window 
- Default Replication Type – **Live Replication (Not Retroactive)** - automatically replicates new and updated objects
- **On-demand replication** – To replicate existing objects using S3 Batch Replication
- **Not bidirectional** by default
- Only **user events** are replicated (not system events)
- **Deletes are not replicated by default** (**Delete Markers Not Replicated**)



# Configure Amazon S3 Replication

Cross Account Replication Configuration



# Amazon S3 Performance

Scaling and Performance Strategies

# Using Prefixes to Improve Performance

Amazon S3 scales well and applications can achieve thousands of requests per second when uploading and retrieving from Amazon S3.

- 3,500 PUT/COPY/POST/DELETE
- 5,500 GET/HEAD
- No limits to the number of prefixes in a bucket
- Increase performance by parallelizing transactions

per prefix in a bucket

## Example

- customerrecipe/*image-a-e*/file...file...file
- customerrecipe/*image-f-j*/file...file...file
- customerrecipe/*image-k-o*/file...file...file
- custoemrrecipe/*image-p-t*/file...file...file
- customerrecipe/*image-u-z*/file...file...file

Spreading reads across all 5 prefixes, and parallelize those reads, gives you 27,500 read requests per second





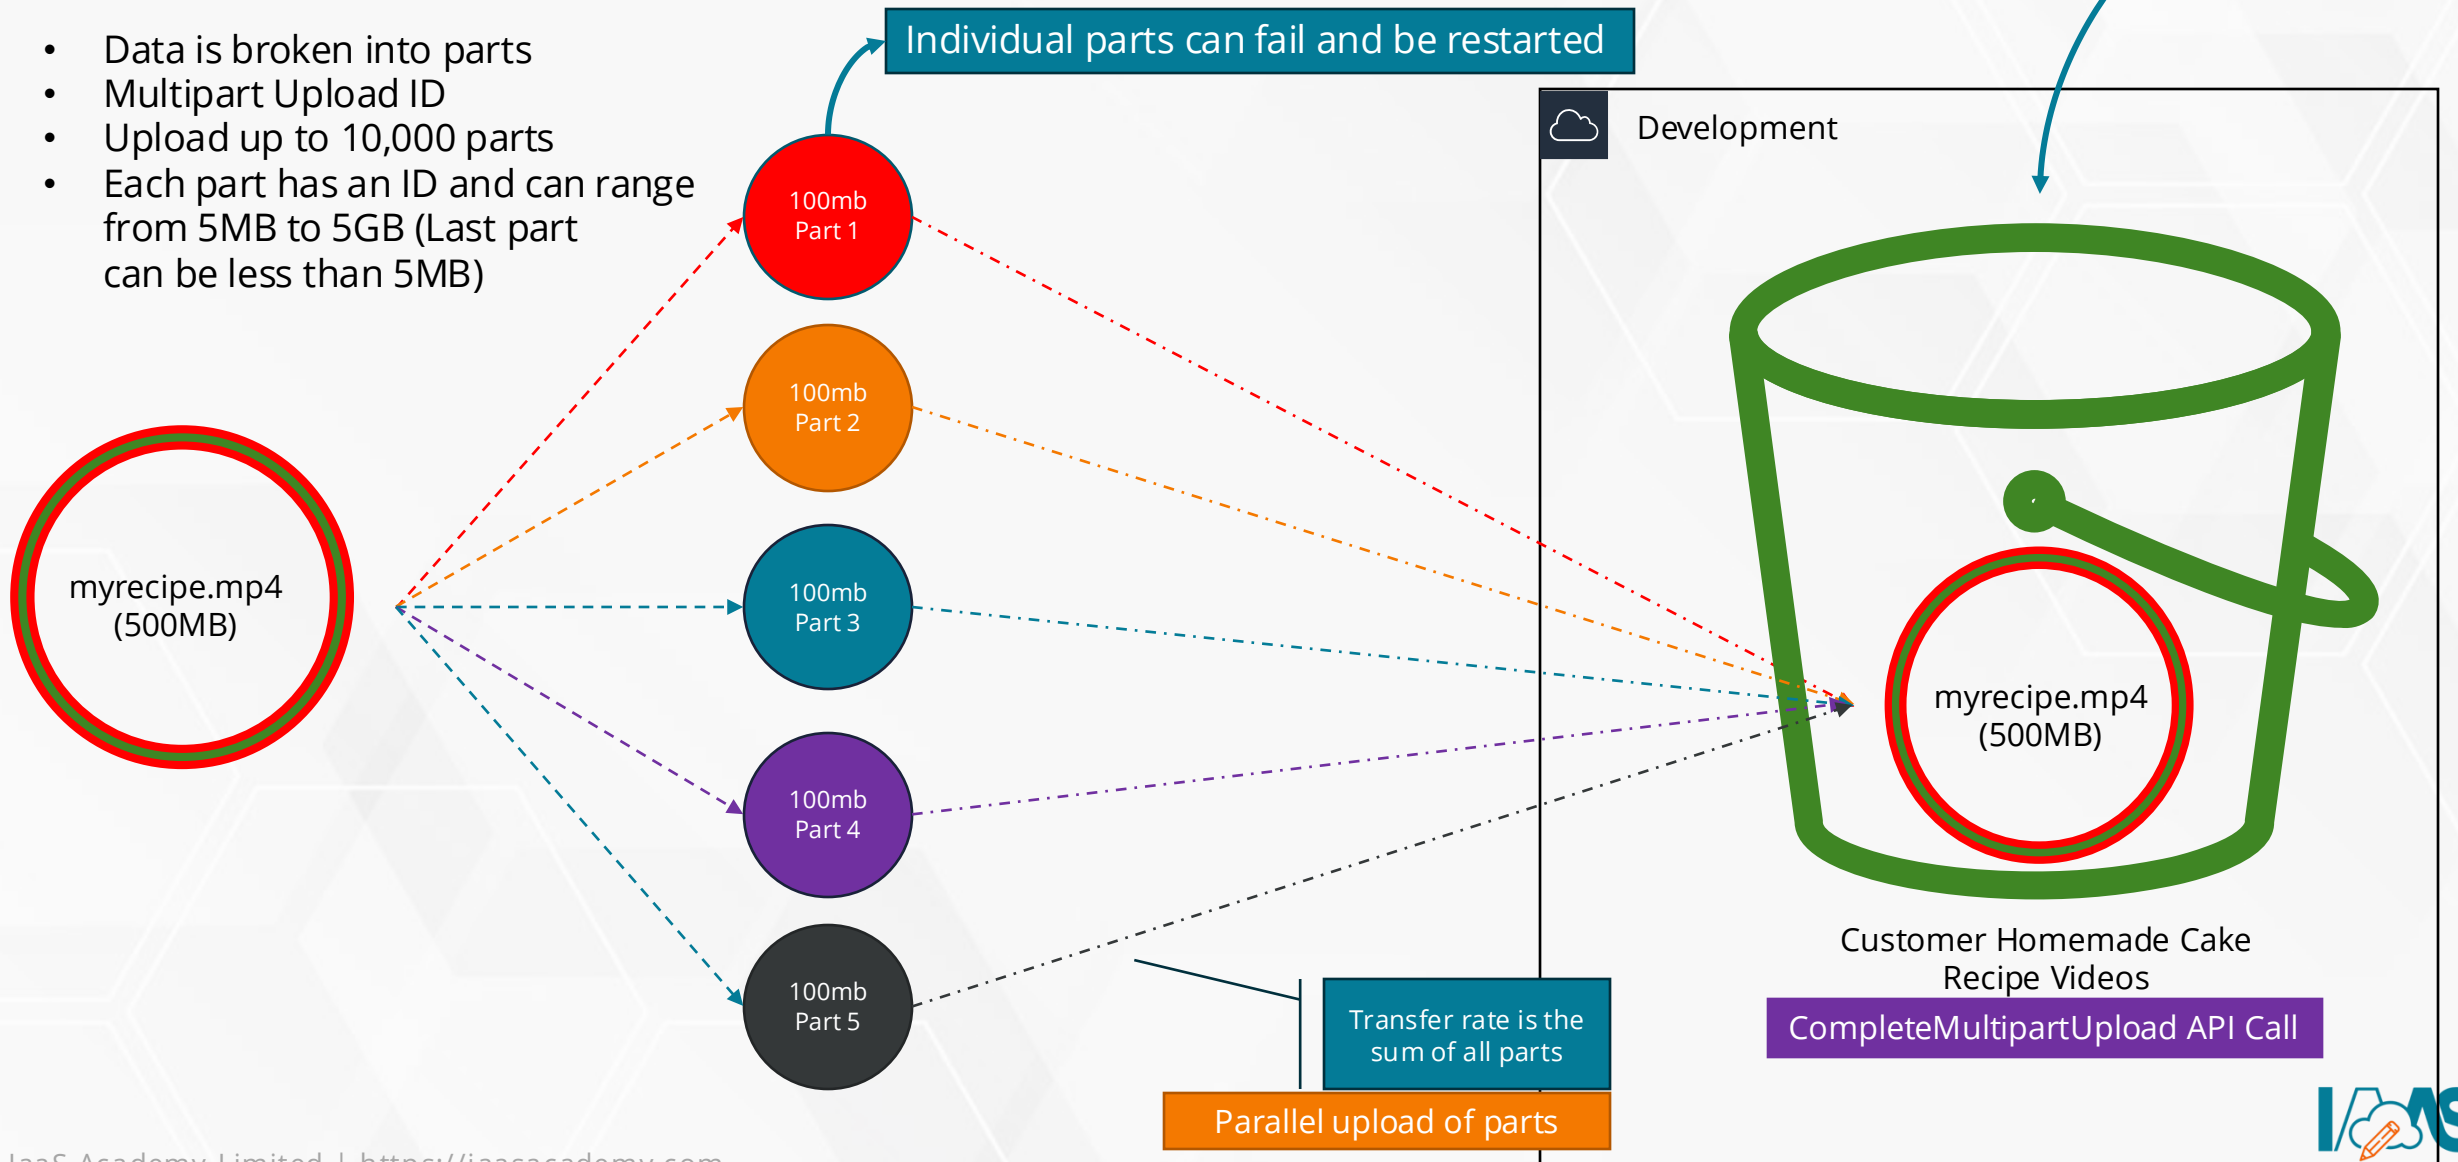
RITUAL  
ROAST

# Customer Scenario – Ritual Roast Vlog Contest

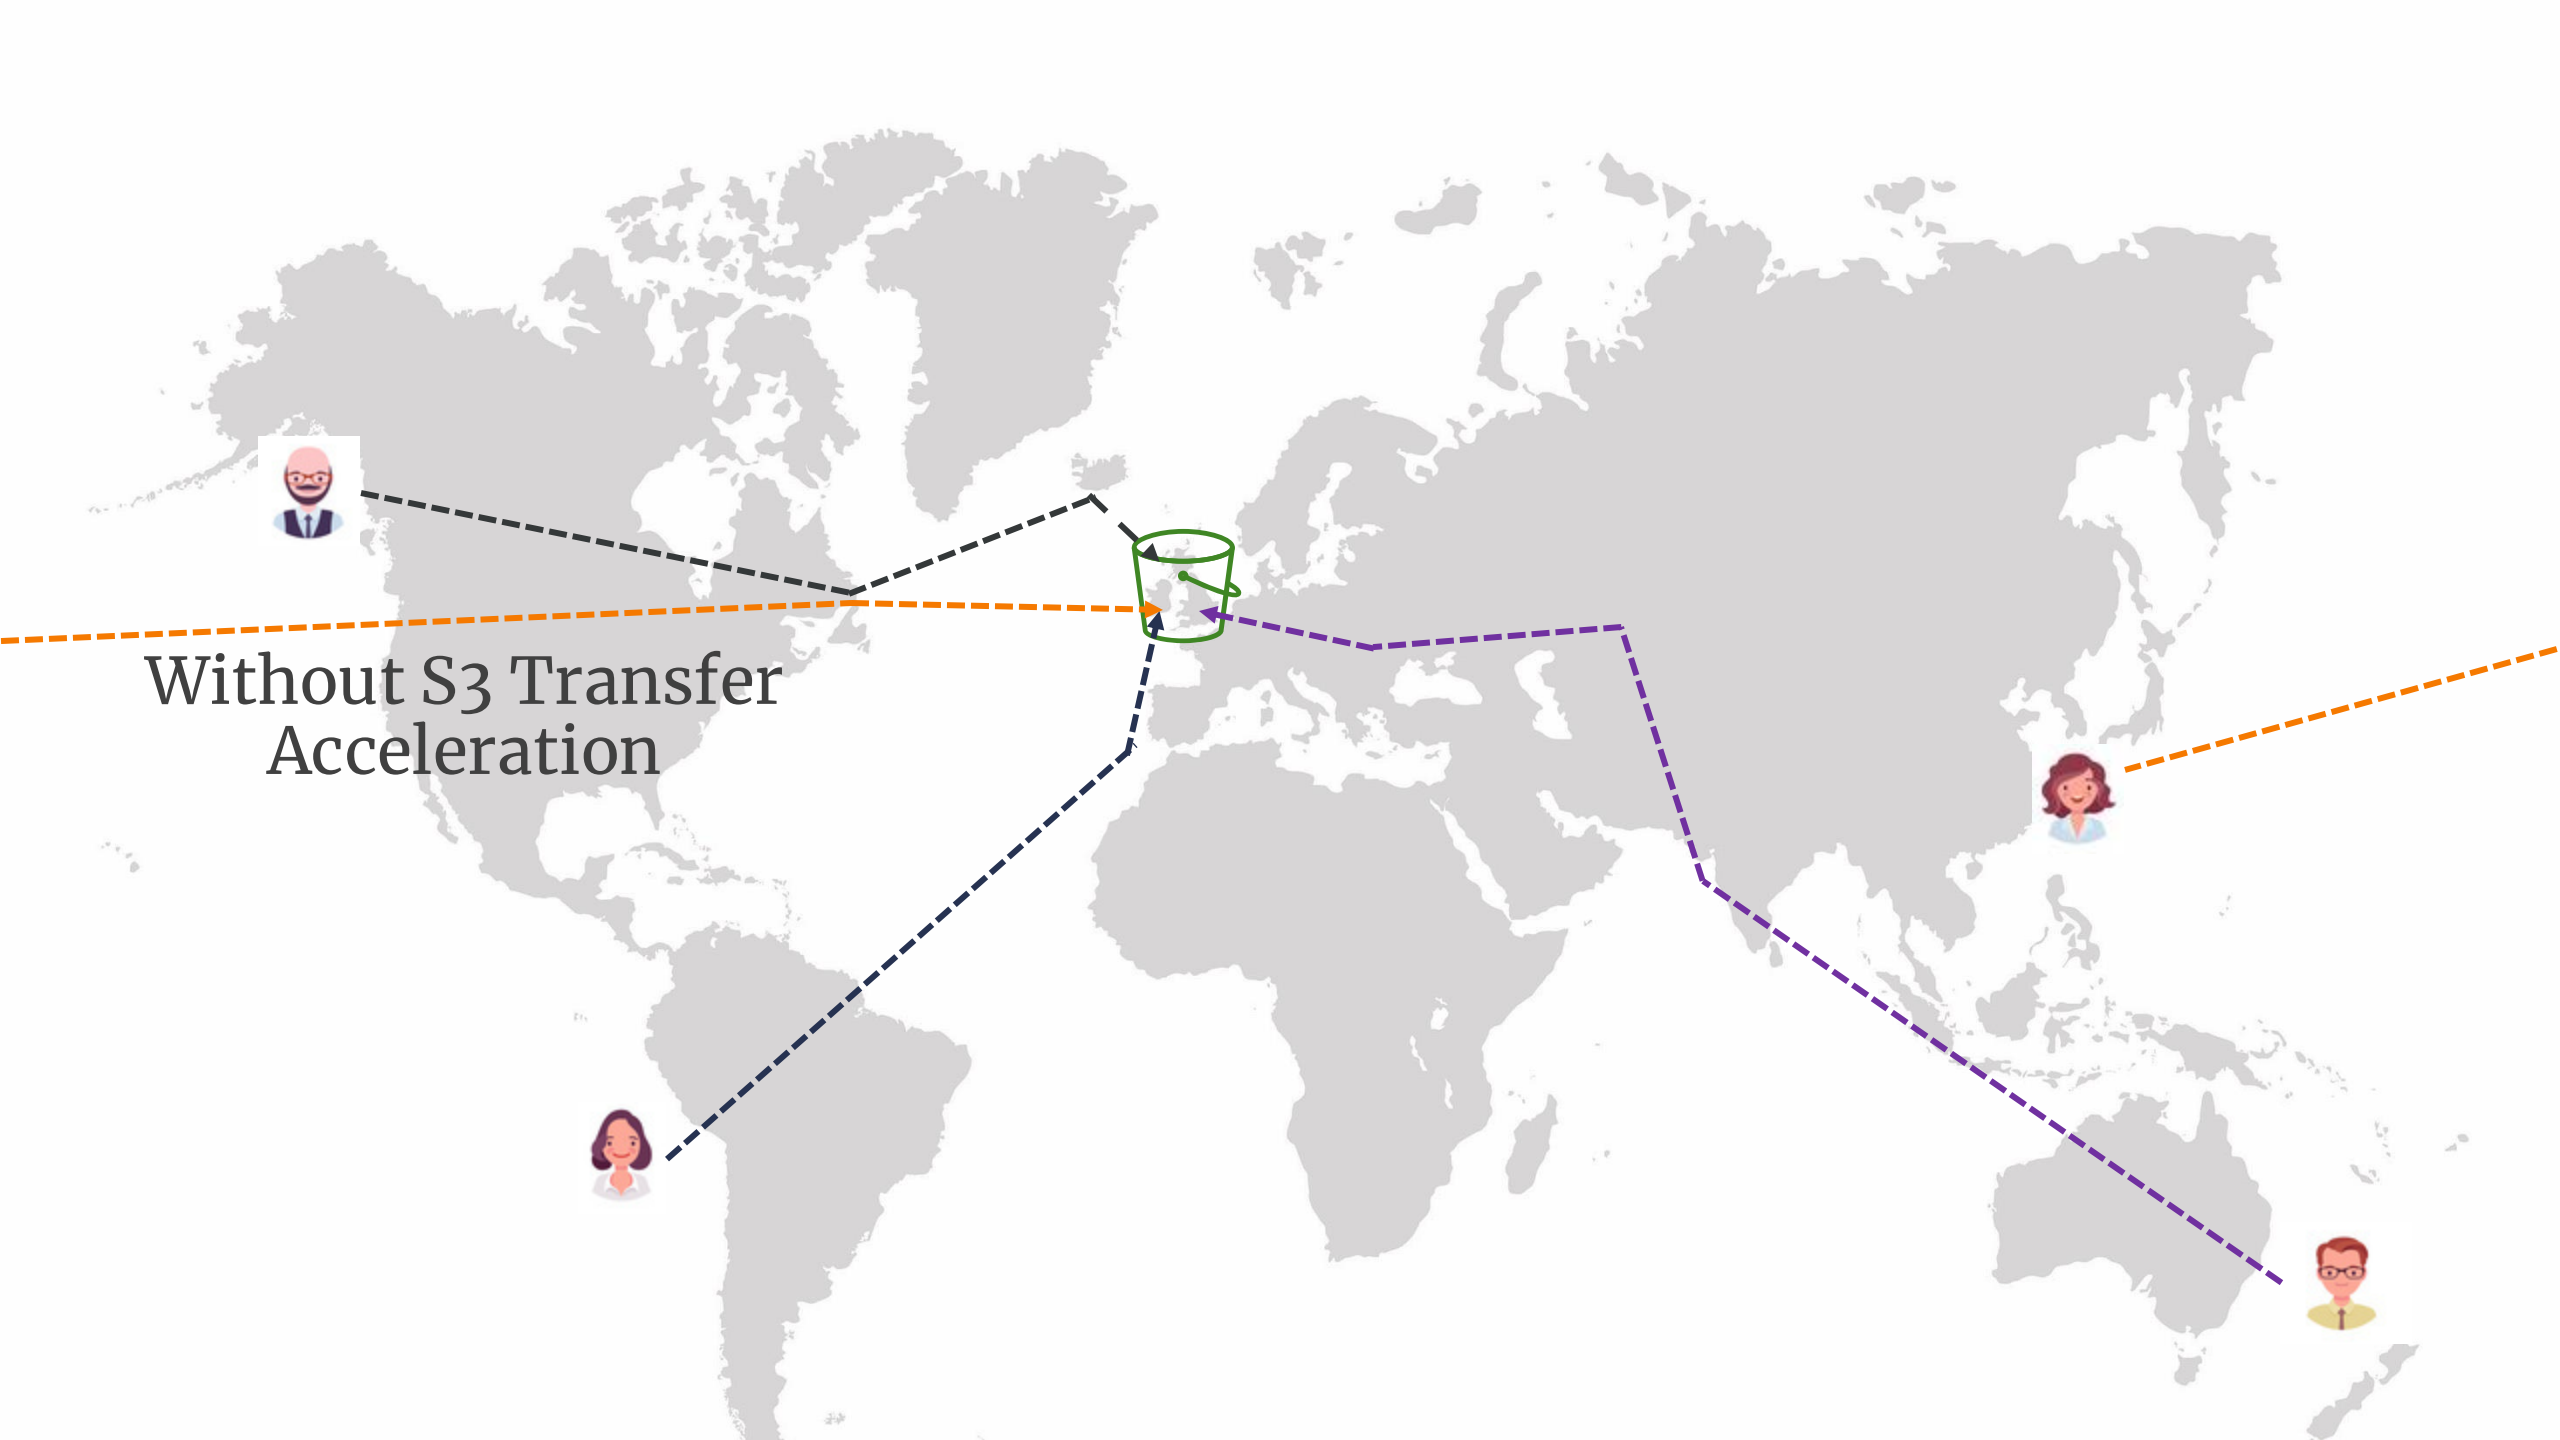


# Multipart Upload

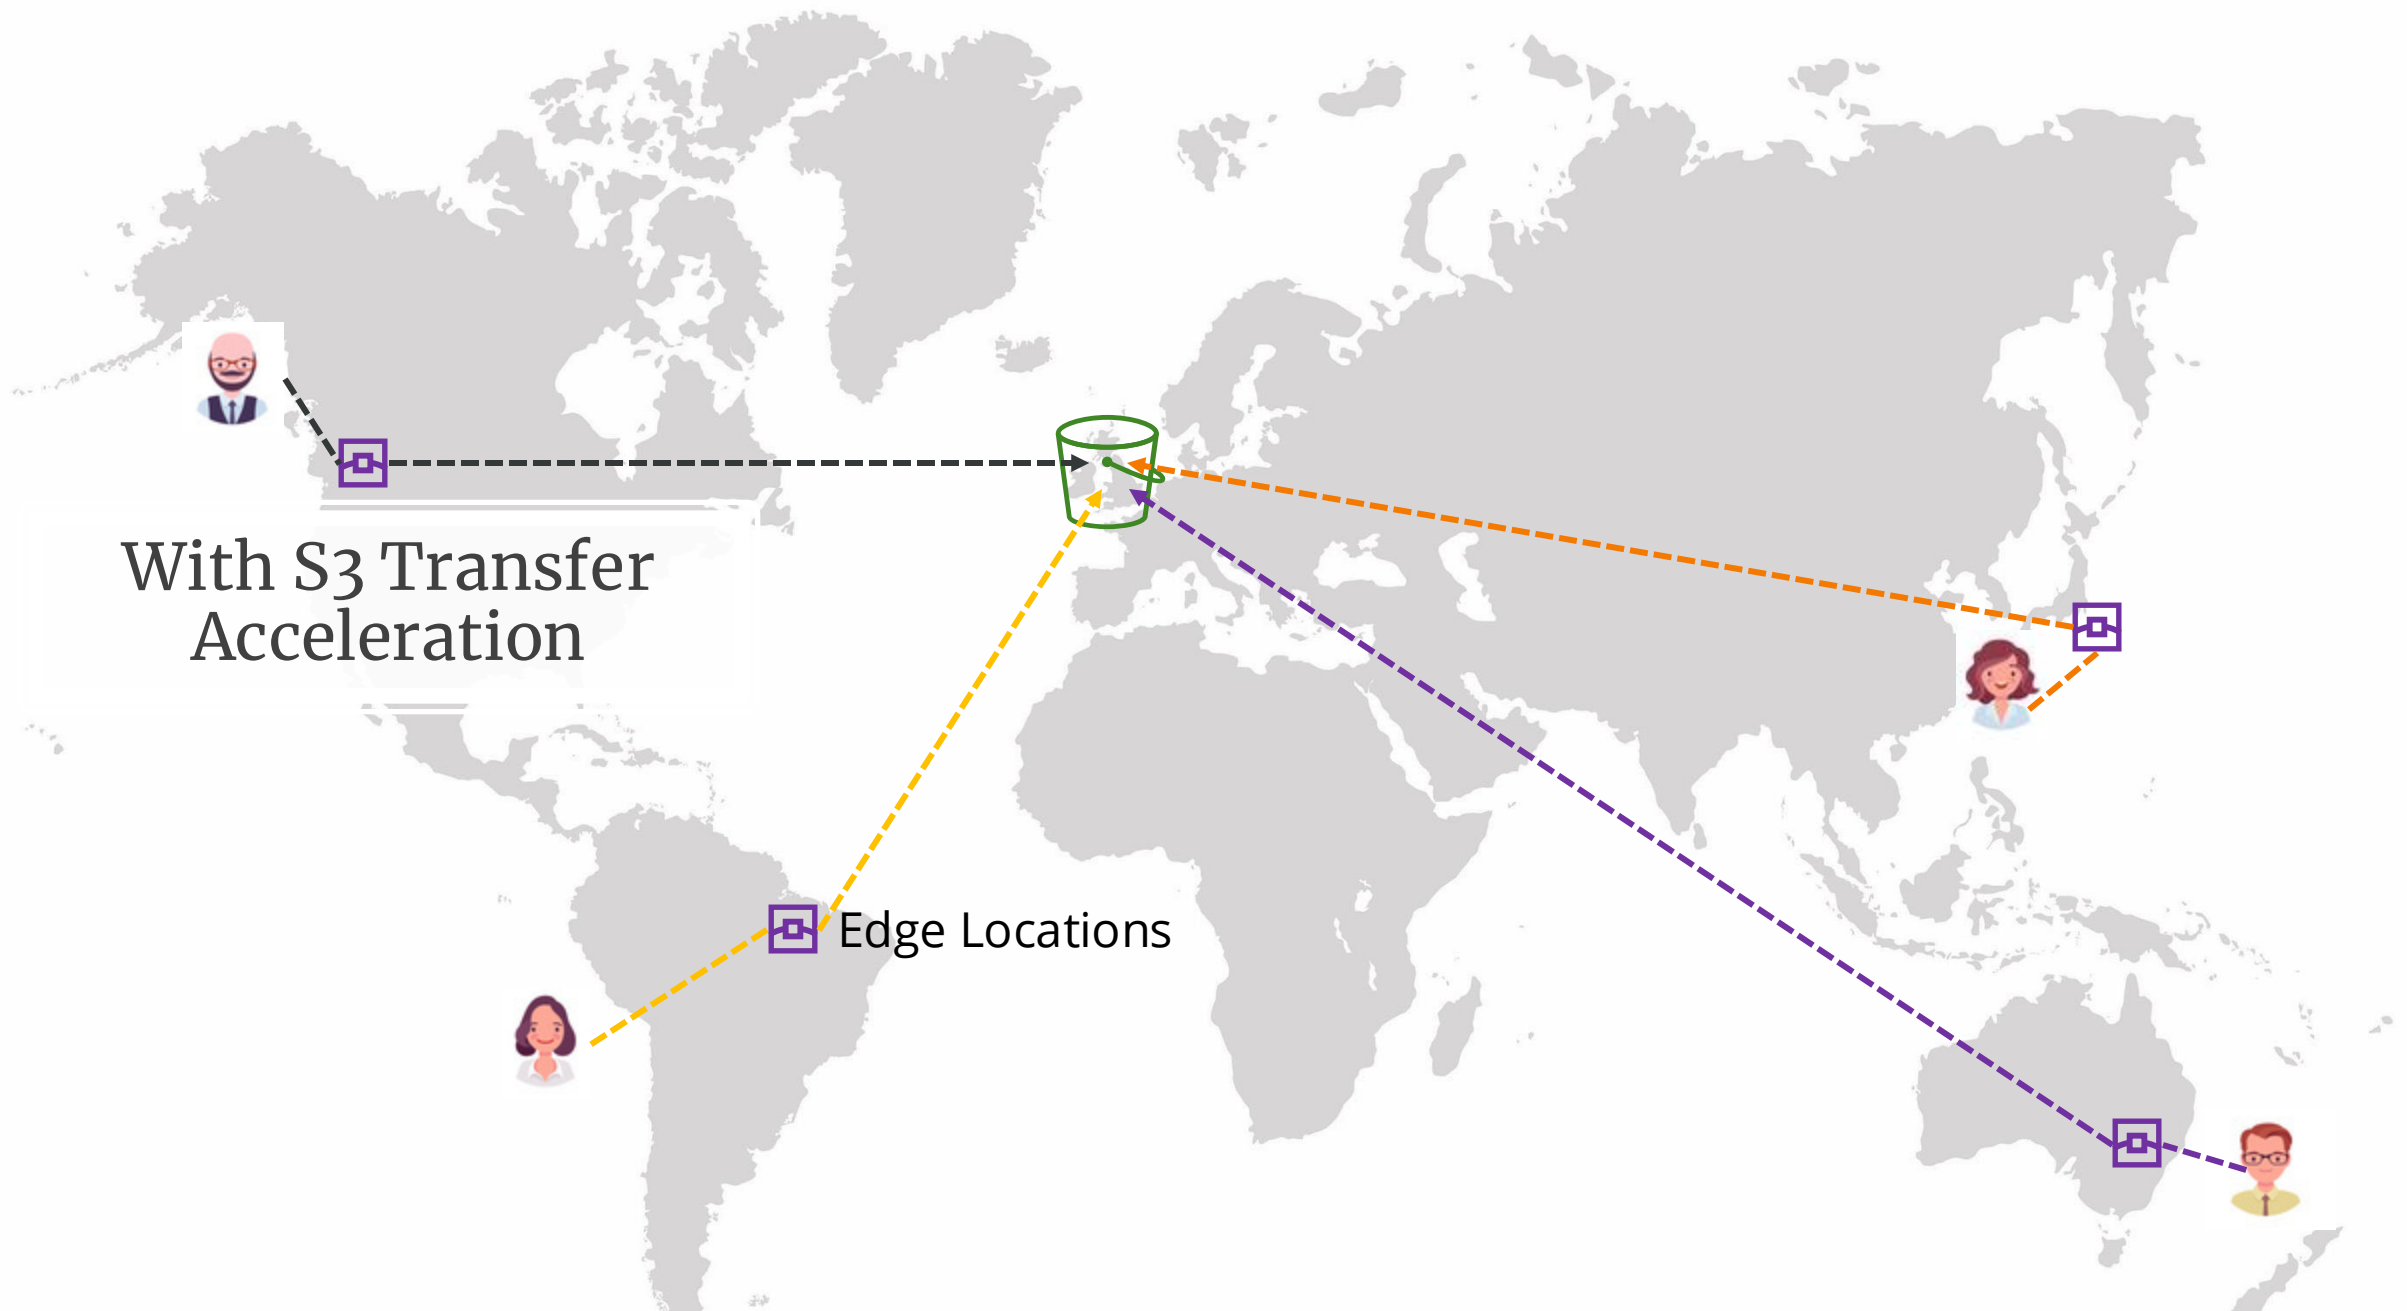
- Data is broken into parts
- Multipart Upload ID
- Upload up to 10,000 parts
- Each part has an ID and can range from 5MB to 5GB (Last part can be less than 5MB)







Without S3 Transfer Acceleration



With S3 Transfer Acceleration

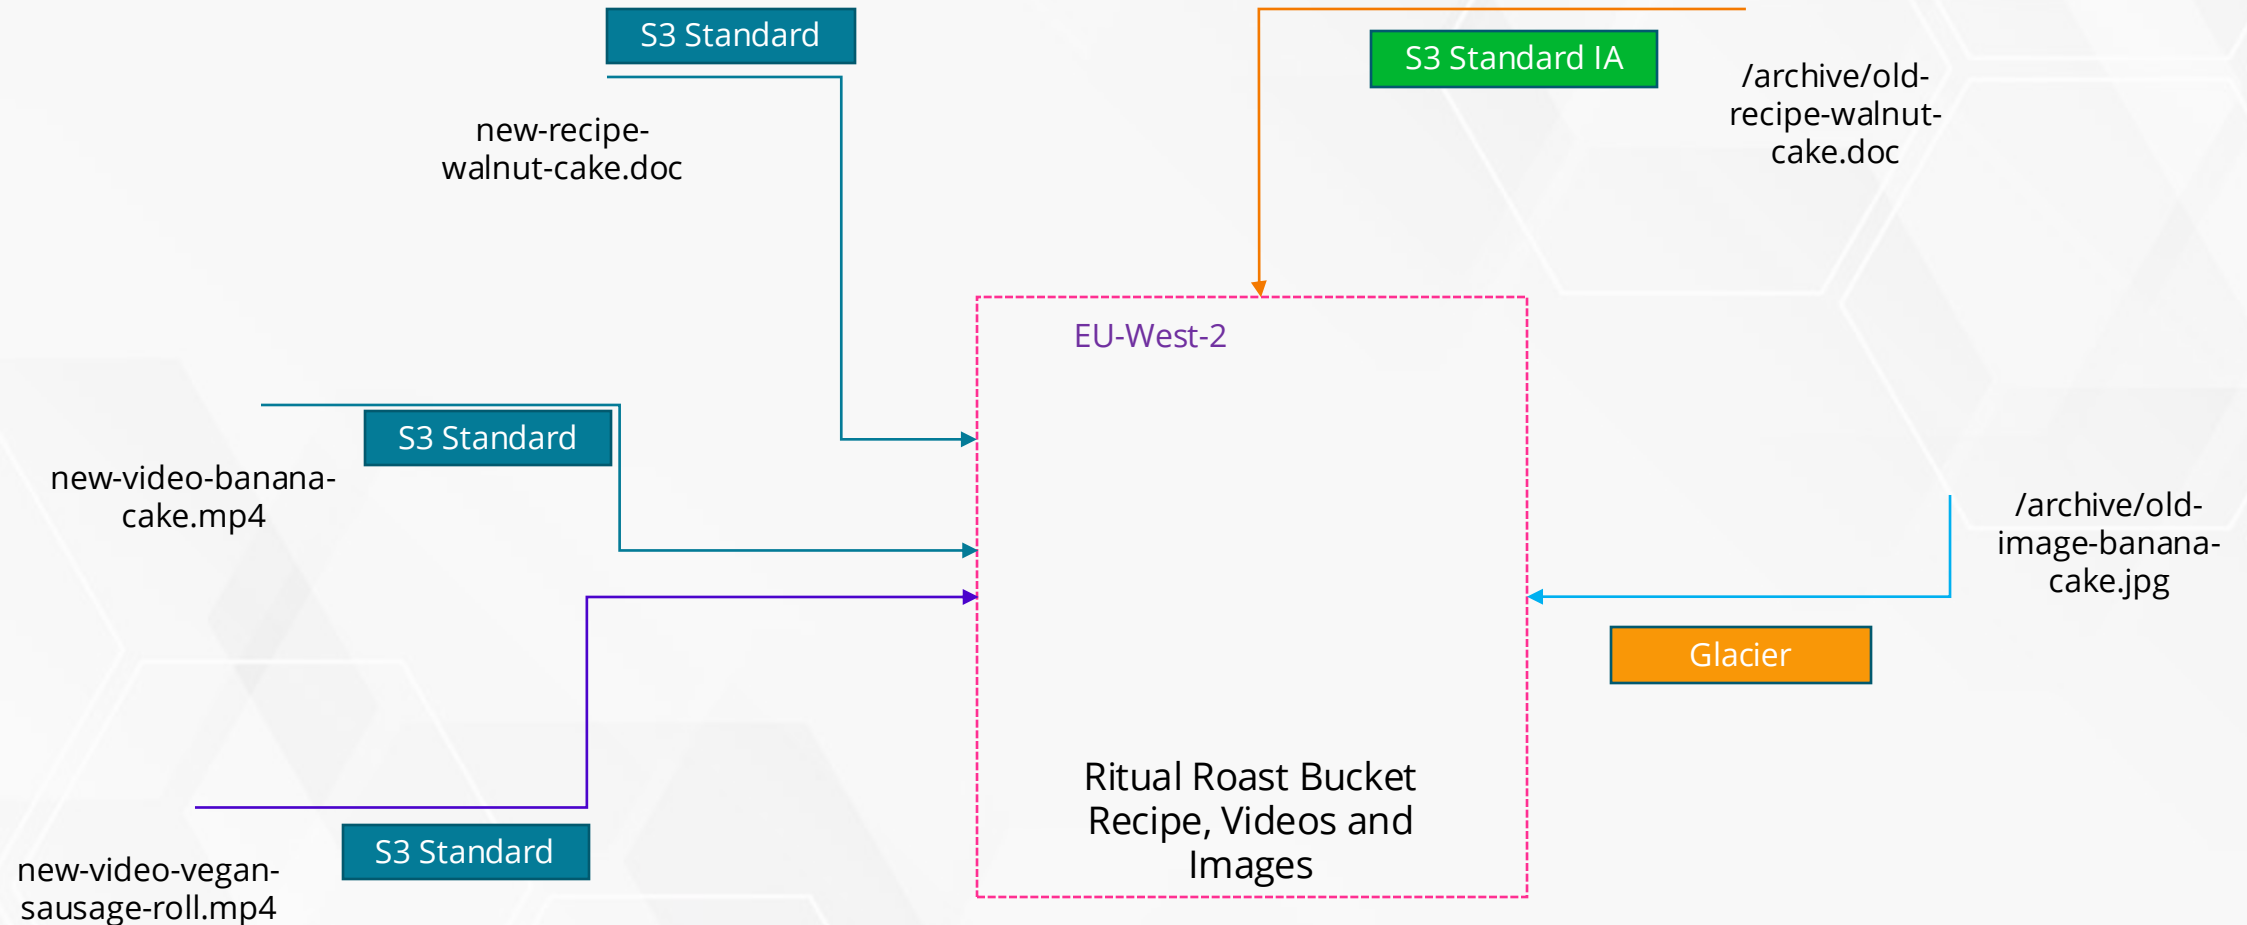
Edge Locations



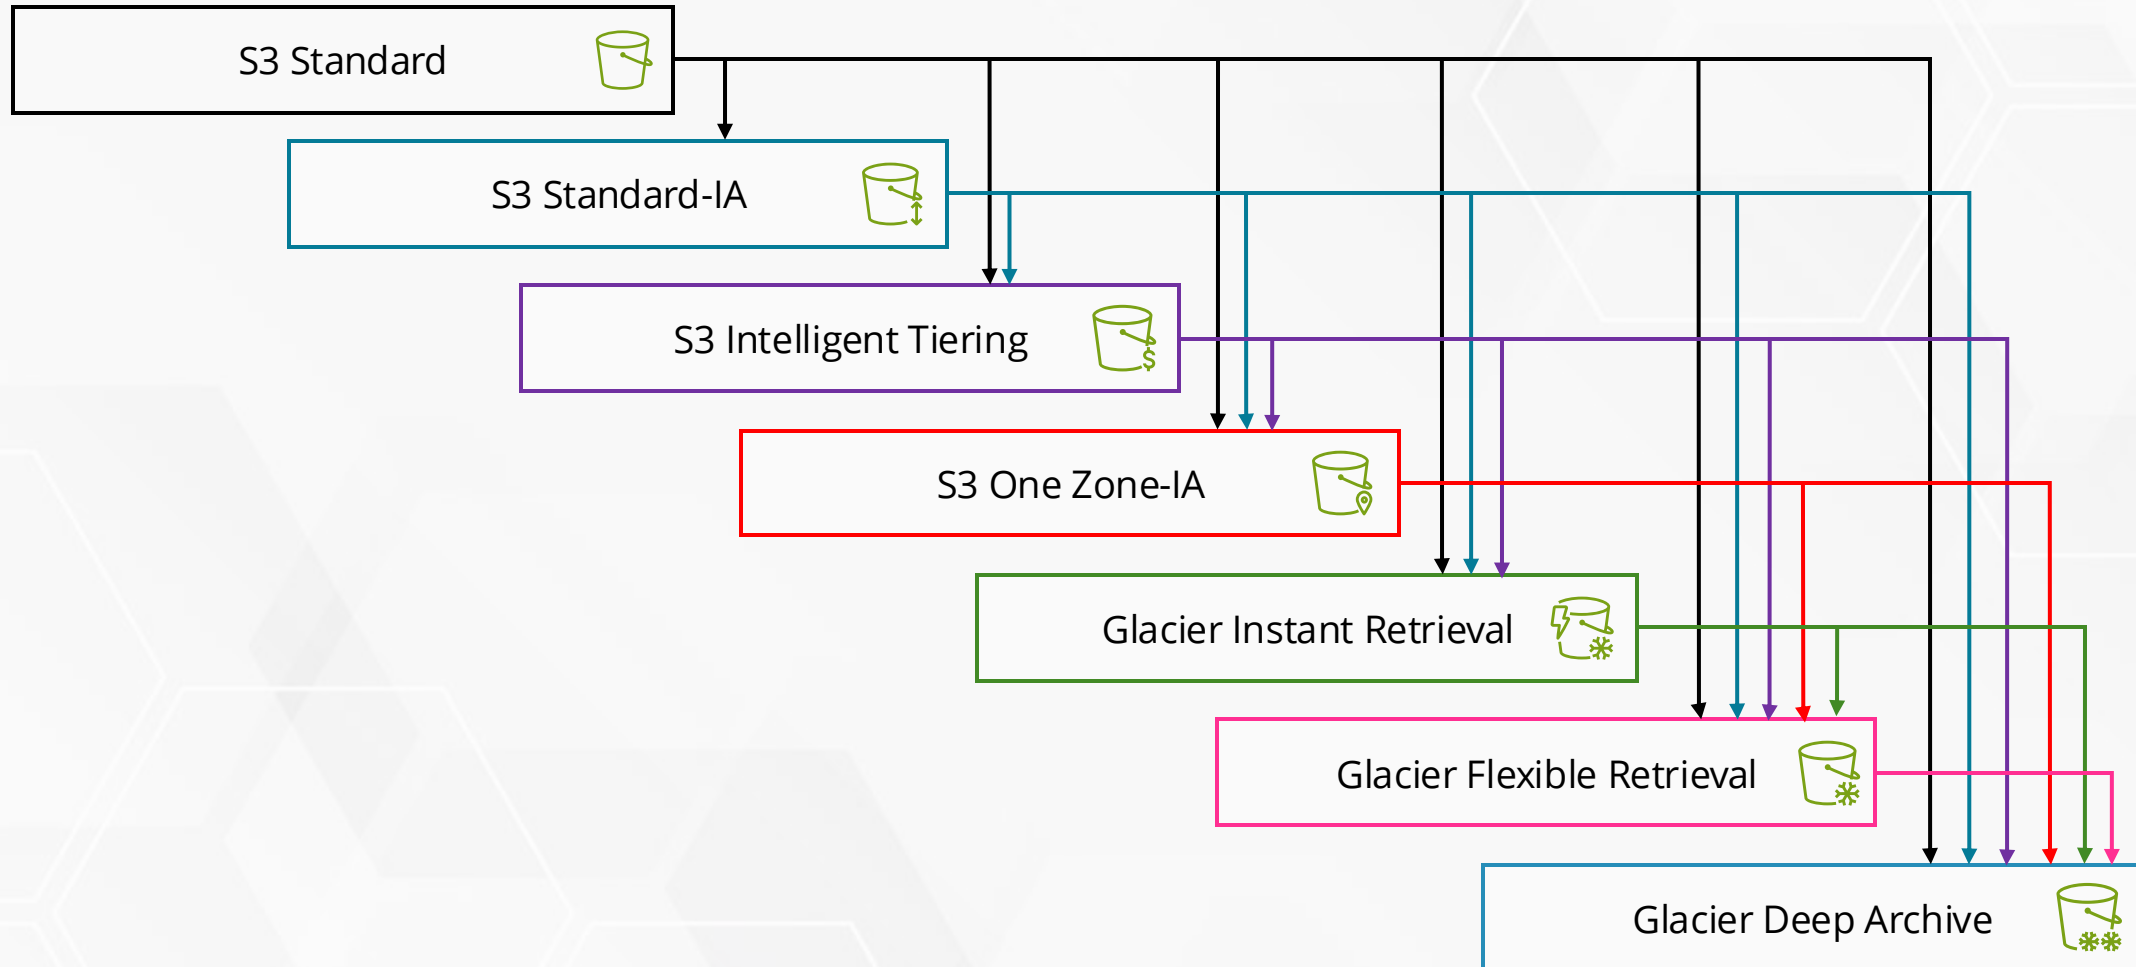
# Amazon S3 Lifecycle Management

Managing Data Lifecycle in your  
Organization

# Amazon S3 Bucket – Objects Storage Classes



# Amazon S3 Lifecycle Management

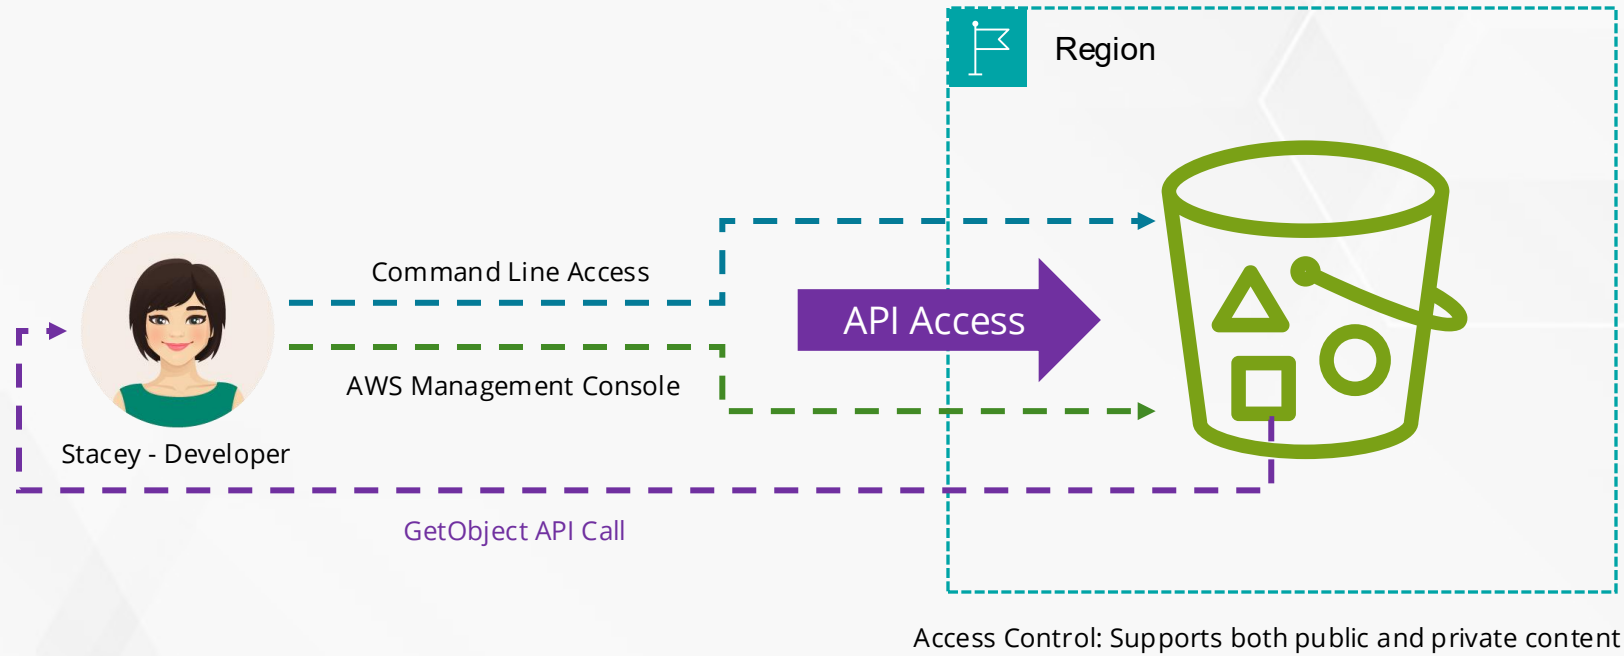




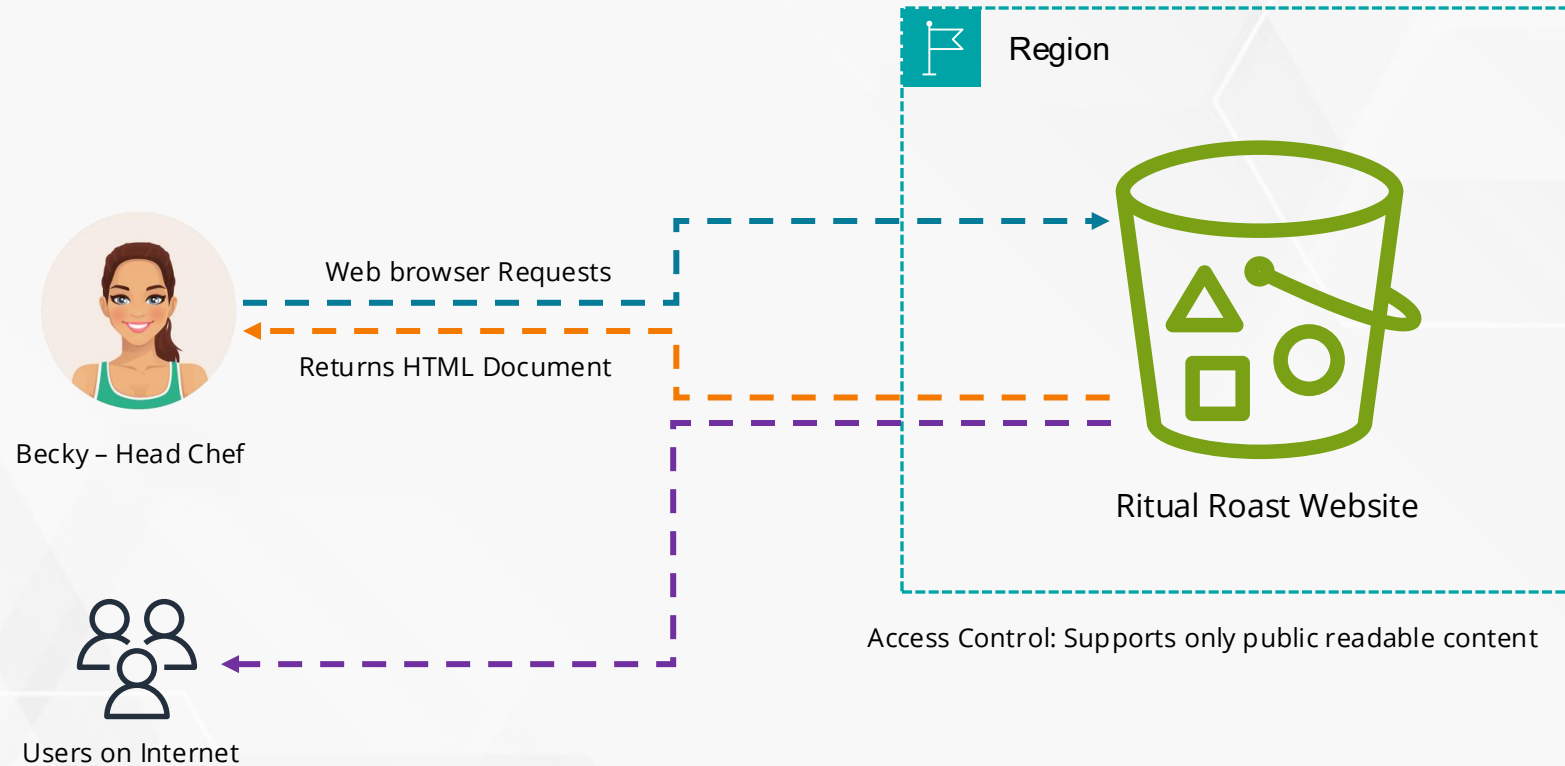
# Amazon S3 Static Website Hosting

Hosting Front-End Static Websites

# Accessing Amazon S3 Buckets – API Access



# Accessing Amazon S3 Buckets – Website Endpoint



## URL Format - Region Dependent

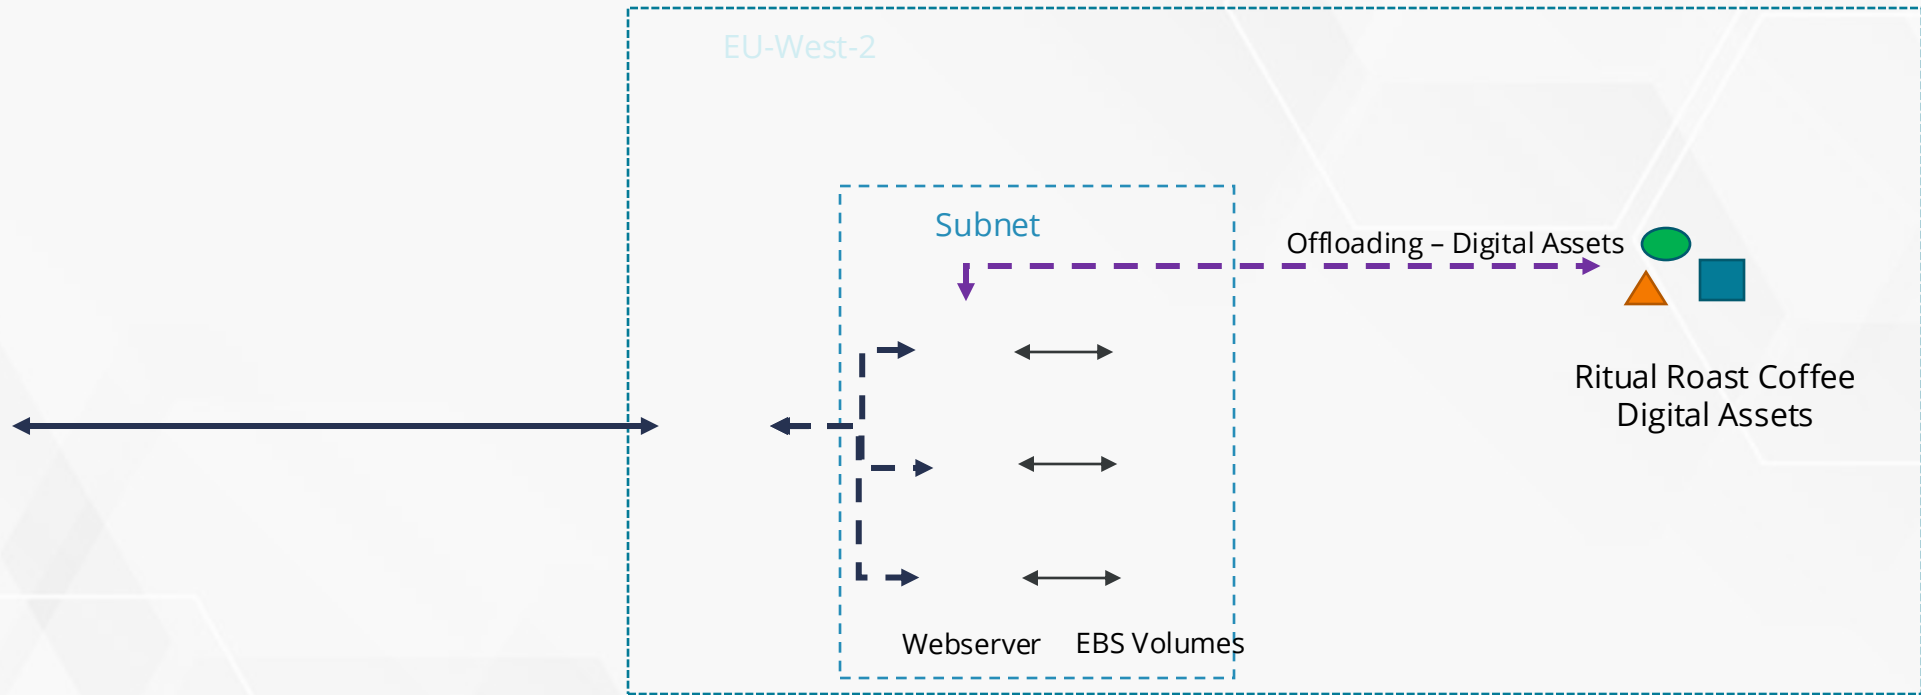
s3-website dash (-) Region - `http://bucket-name.s3-website.Region.amazonaws.com`  
s3-website dot (.) Region - `http://bucket-name.s3-website-Region.amazonaws.com`




# Static Website Hosting – Product Launch



# Static Website Hosting – Offloading



# Static Website Hosting – Under Maintenance



# Static Website Hosting

To configure a bucket to host a static website, follow these steps:

- Create a bucket with the name that will be your website hostname
- Upload static files to the bucket
  - You will need an **index.html** file
  - You will need an **error.html** file
- Make all files public or grant everyone read access
- Enable the bucket for Static Website Hosting
- Optionally configure a DNS CNAME to point to the Amazon URL. However, your **bucket name must match the custom domain name.**

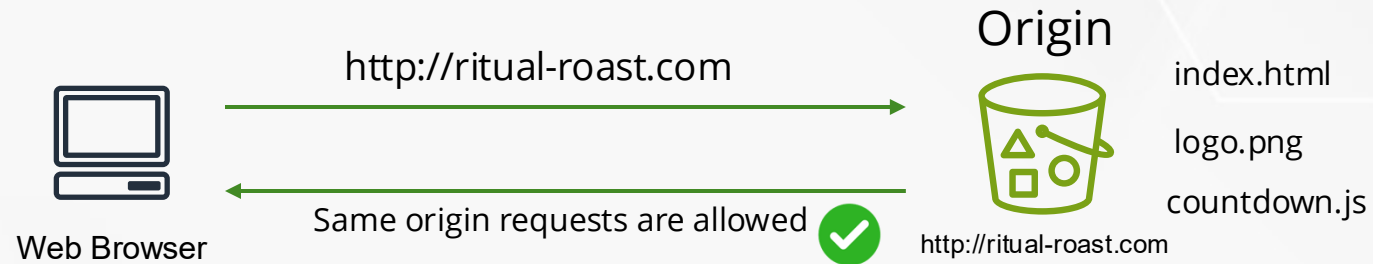


# Cross Origin Resource Sharing

Introduction to CORS

# What is CORS

Browser security feature allows client web applications loaded in one domain to interact with resources in a different domain.

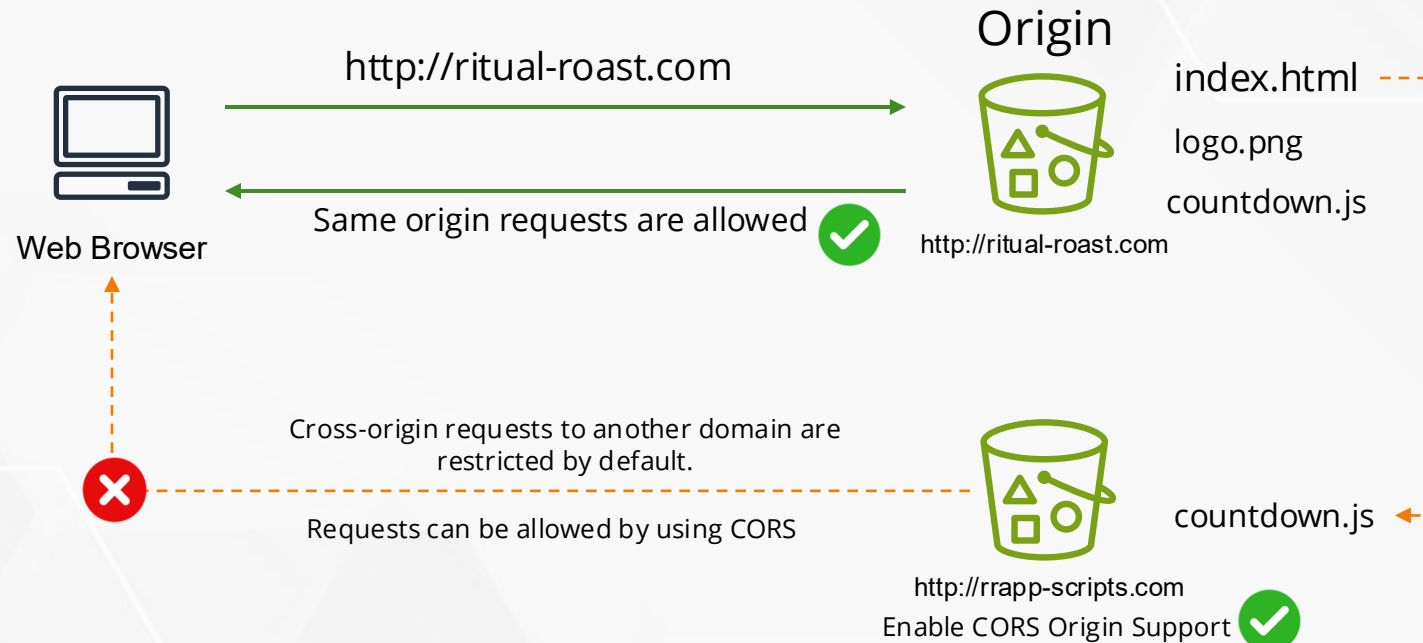


# What is CORS

Browser security feature allows client web applications loaded in one domain to interact with resources in a different domain.

Origin = scheme(protocol) + host(host) + port

Example: <http://ritual-roast.com> (port 80 for HTTP or 443 for HTTPS)



# Additional Points

---

- Simple Requests
- Preflight requests & responses
- Access-Control-Allow-Origin (specify the origins that you want to allow cross-domain requests from)
- Access-Control-Max-Age
- Access-Control-Allow-Methods (GET, PUT, POST, DELETE, HEAD)
- Access-Control-Allow-Headers



# Example CORS Configuration

```
[
  {
    "AllowedHeaders": [
      "*"
    ],
    "AllowedMethods": [
      "PUT",
      "POST",
      "DELETE"
    ],
    "AllowedOrigins": [
      "http://www.example1.com"
    ],
    "ExposeHeaders": []
  },
  {
    "AllowedHeaders": [],
    "AllowedMethods": [
      "GET"
    ],
    "AllowedOrigins": [
      "*"
    ],
    "ExposeHeaders": []
  }
]
```

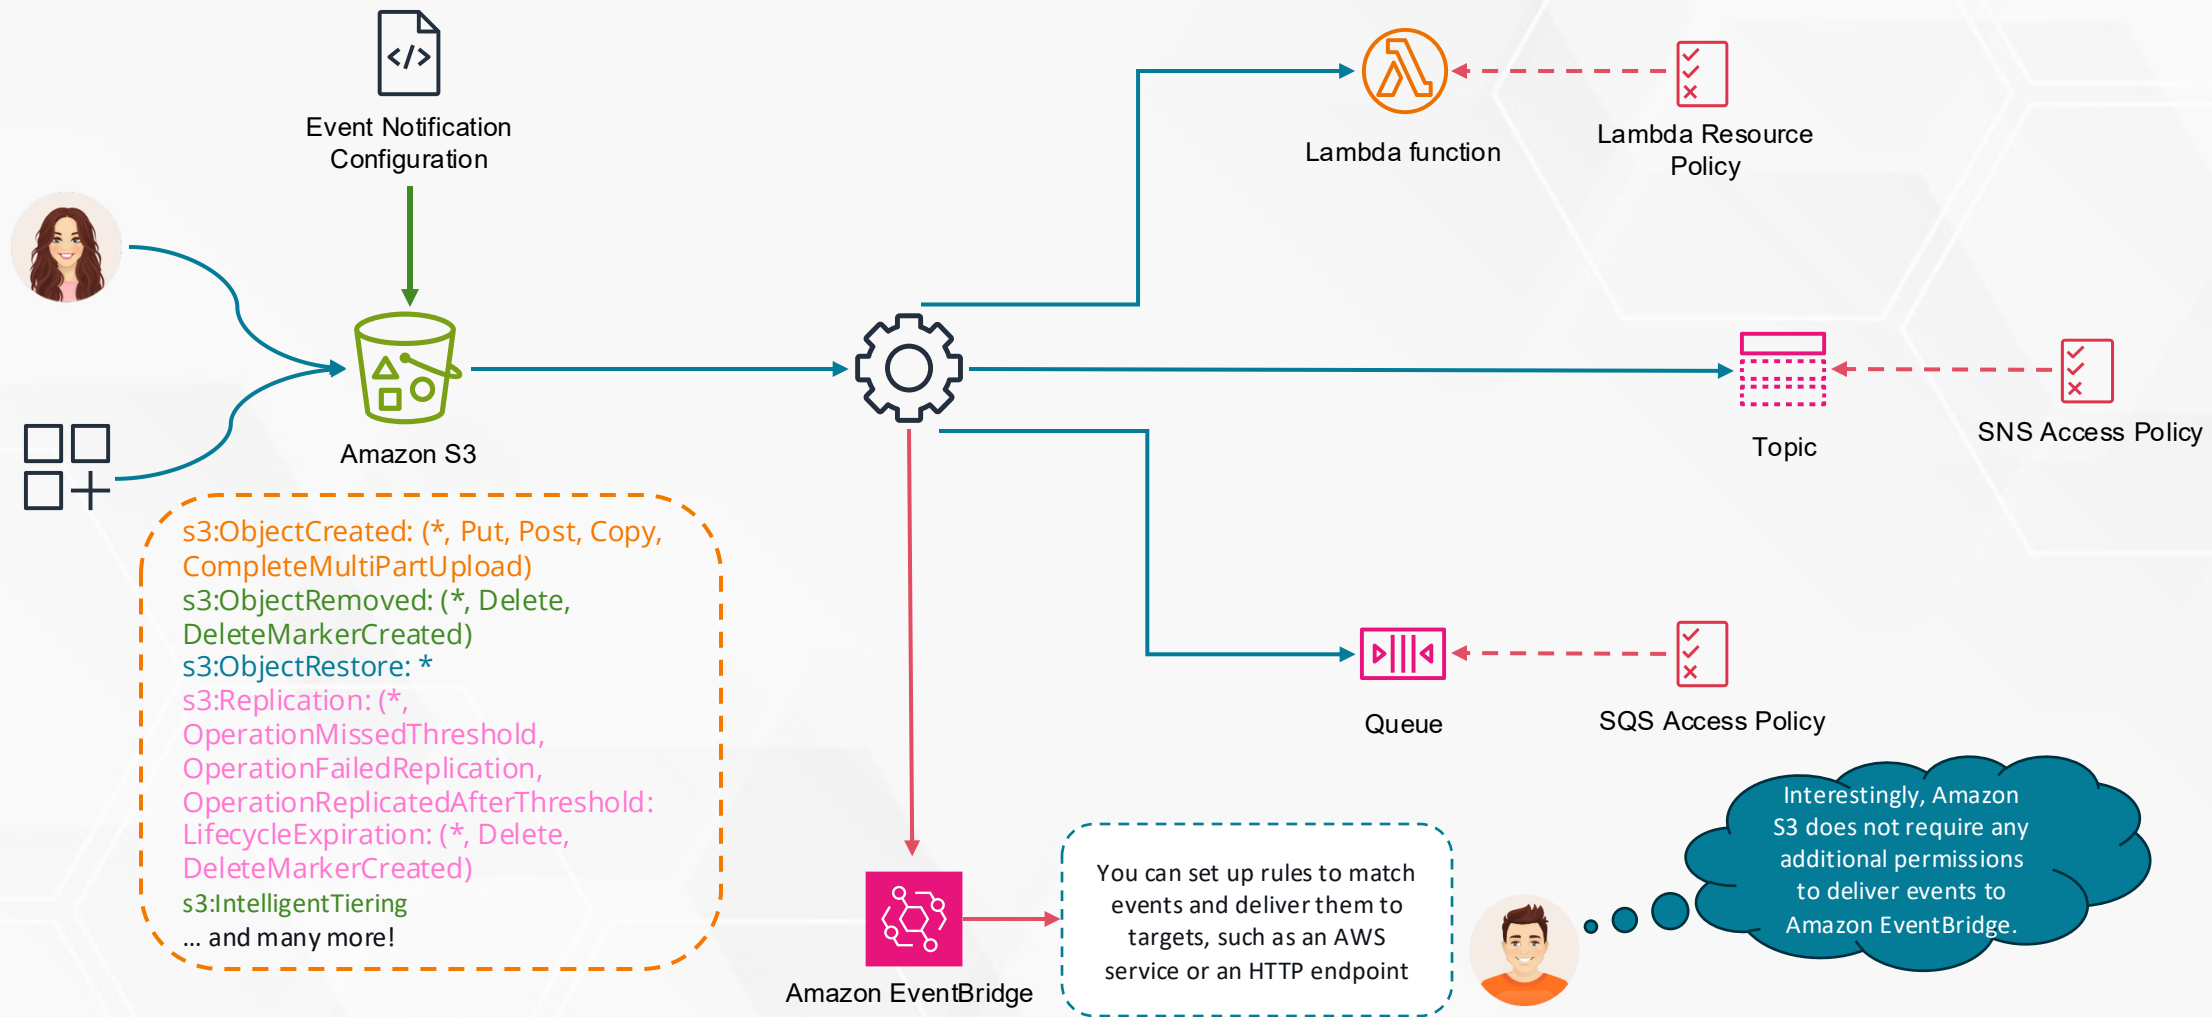




# Amazon S3 Event Notification

# Amazon S3 Event Notification

Use the Amazon S3 Event Notifications feature to receive notifications when certain events happen in your S3 bucket



# Sample SNS Access Policy



```
{
  "Version": "2012-10-17",
  "Id": "example-ID",
  "Statement": [
    {
      "Sid": "Example SNS topic policy",
      "Effect": "Allow",
      "Principal": {
        "Service": "s3.amazonaws.com"
      },
      "Action": [
        "SNS:Publish"
      ],
      "Resource": "SNS-topic-ARN",
      "Condition": {
        "ArnLike": {
          "aws:SourceArn": "arn:aws:s3:*:*:amzn-s3-demo-bucket"
        },
        "StringEquals": {
          "aws:SourceAccount": "bucket-owner-account-id"
        }
      }
    }
  ]
}
```



# Amazon S3 Encryption Options

Encrypting Your Amazon S3 Objects

# Introduction to Amazon S3 Encryption

## Encryption **in transit** vs **at rest**



With Client-side encryption:

- Your objects aren't exposed to any third party, including AWS
- Amazon S3 does not play a role in encrypting or decrypting your objects
- You must use the Amazon S3 Encryption Client to automatically encrypt and decrypt as part of your Amazon S3 PutObject and GetObject requests

Client-side Encryption



HTTPS TUNNEL



Amazon S3

£\$%^@\*&!(!@£



Server-side Encryption



HTTPS TUNNEL



Amazon S3

PLAIN TEXT



£%\$&@%^£



Encryption is defined at a per object level, not bucket level. Each object might use a different encryption setting

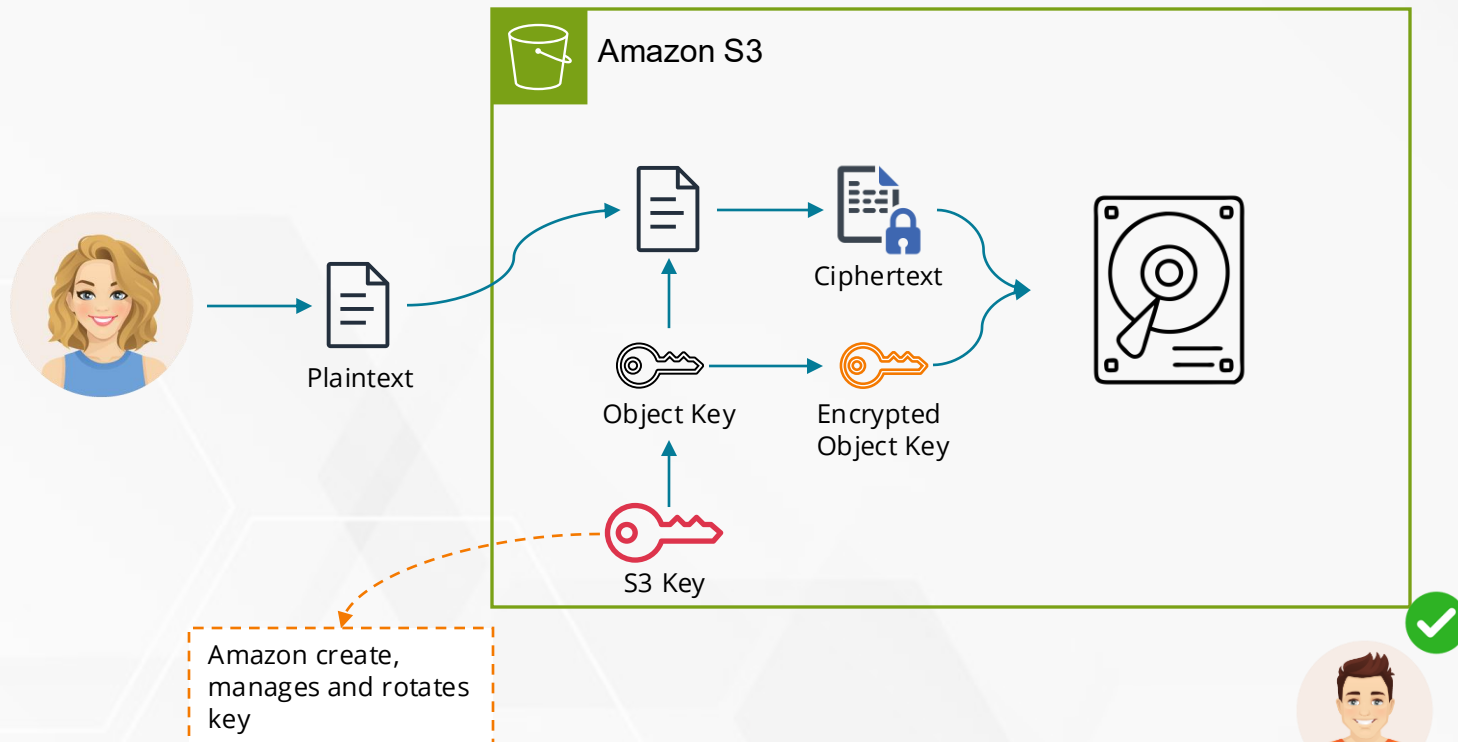
Server-side encryption is now mandatory

# Introduction to Amazon S3 Encryption

AWS manages creation, deletion and usage of keys – Zero Control on Keys!

Other users with admin access can access data = **No Role Separation**

## Server-side encryption with Amazon S3 managed keys (SSE-S3)



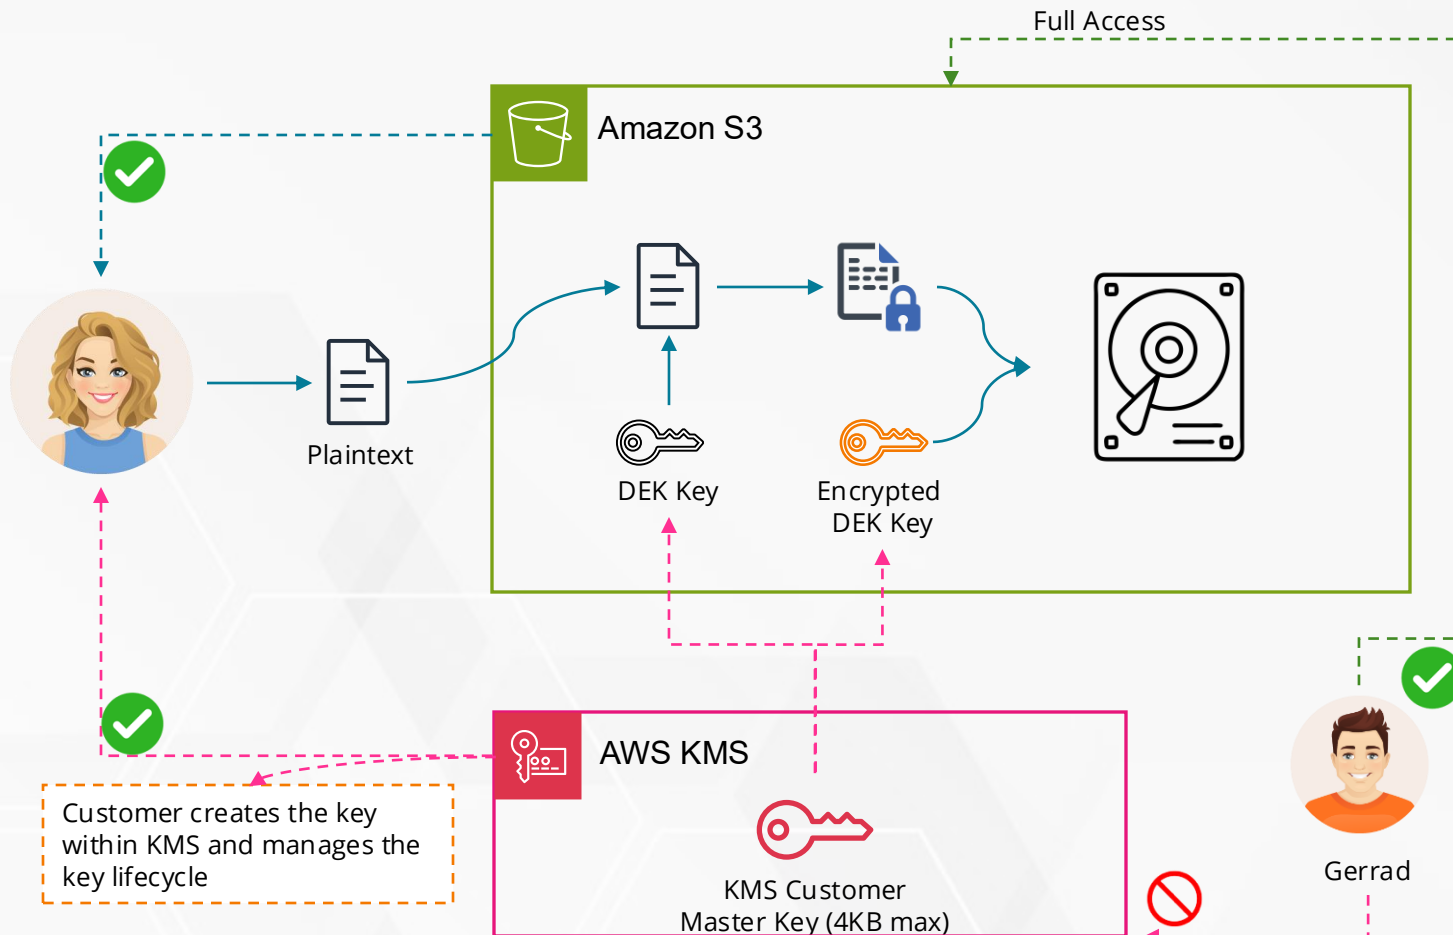
- **Server-side encryption with Amazon S3 managed keys (SSE-S3)** – Default encryption (no cost)
- AWS creates and manages the keys for you
- SSE-S3 is applied to all objects uploaded to the bucket
- Using object creation REST APIs, you must provide the **"x-amz-server-side-encryption": "AES256"** request header

Gerrad  
S3 Administrator  
Full Access

# Introduction to Amazon S3 Encryption

Note: AWS KMS keys must be in the same Region as the bucket.

## Server-side encryption with AWS KMS keys (SSE-KMS)



Customer creates the key within KMS and manages the key lifecycle

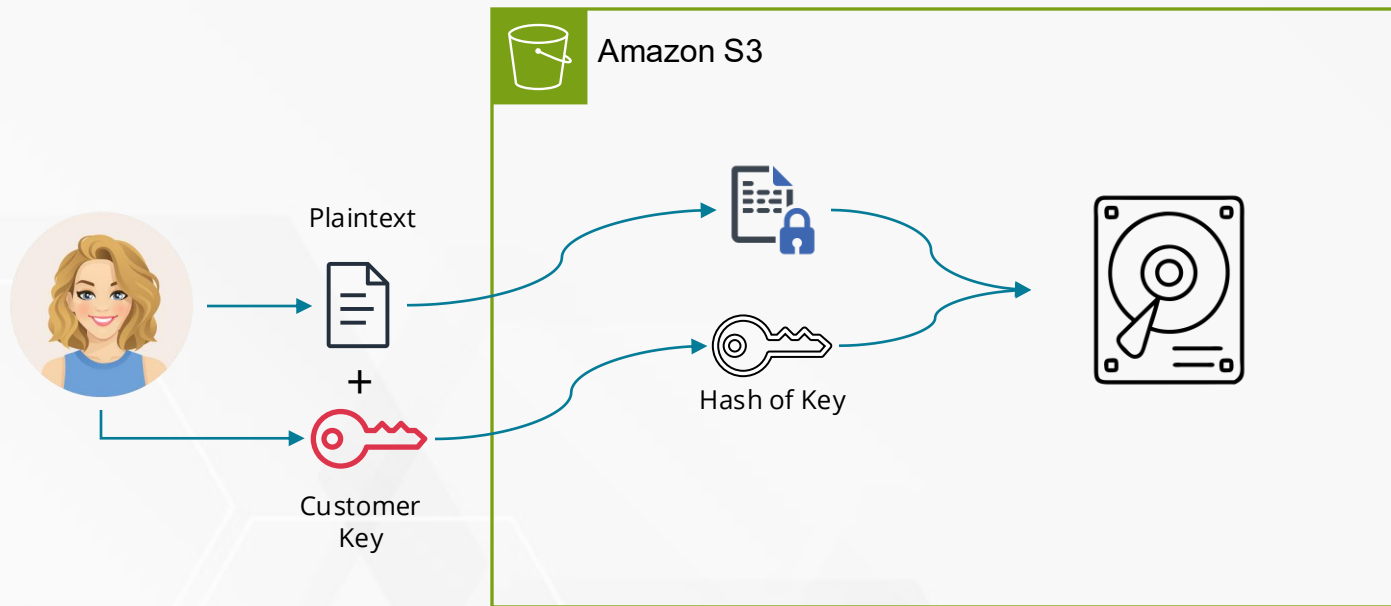
- **Server-side encryption with AWS KMS keys (SSE-KMS)**
- Centrally create, view, edit, monitor, enable or disable, rotate, and schedule deletion of KMS keys.
- Define the policies that control how and by whom KMS keys can be used – **Role Separation**
- Audit their usage to prove that they are being used correctly (CloudTrail)
- Ideal for regulated industries
- Using object creation REST APIs, you must provide the **"x-amz-server-side-encryption": "aws:kms"** request header.

Dual-layer server-side encryption with AWS Key Management Service (AWS KMS) keys (DSSE-KMS) applies two layers of encryption to objects when they are uploaded to Amazon S3.



# Introduction to Amazon S3 Encryption

## Server-side encryption with Customer Provided Keys(SSE-C)



Amazon S3 stores a randomly salted Hash-based Message Authentication Code (HMAC) value of the encryption key to validate future requests. The HMAC cannot be used to derive the value of the encryption key or to decrypt the contents of the encrypted object. That means if you lose the encryption key, you lose the object.

- **Server-side encryption with Customer Provided Keys (SSE-C)**
- You store your data encrypted with your own encryption keys
- Amazon S3 manages data encryption as it writes to disks and data decryption when you access your objects.
- Amazon S3 then removes the encryption key from memory
- To retrieve an object, you must provide the same encryption key as part of your request.



# Amazon S3 Encryption – KMS

Additional Concepts on KMS

# Encrypting Existing Objects

- All new buckets have default SSE-S3 encryption enforced

## Default encryption [Info](#)

Server-side encryption is automatically applied to new objects stored in this bucket.

### Encryption type [Info](#)

- ☒ Server-side encryption with Amazon S3 managed keys (SSE-S3)
- ☐ Server-side encryption with AWS Key Management Service keys (SSE-KMS)
- ☐ Dual-layer server-side encryption with AWS Key Management Service keys (DSSE-KMS)

Secure your objects with two separate layers of encryption. For details on pricing, see [DSSE-KMS pricing](#) on the Storage tab of the [Amazon S3 pricing page](#). [↗](#)

### Bucket Key

Using an S3 Bucket Key for SSE-KMS reduces encryption costs by lowering calls to AWS KMS. S3 Bucket Keys aren't supported for DSSE-KMS. [Learn more](#) [↗](#)

- ☐ Disable
- ☒ Enable

What about existing objects in older buckets. How do we encrypt those objects?



Menio

That's easy, you can use AWS S3 Batch Operations to encrypt existing objects. You can encrypt existing objects using **CopyObject** API operation or **copy-object** CLI command

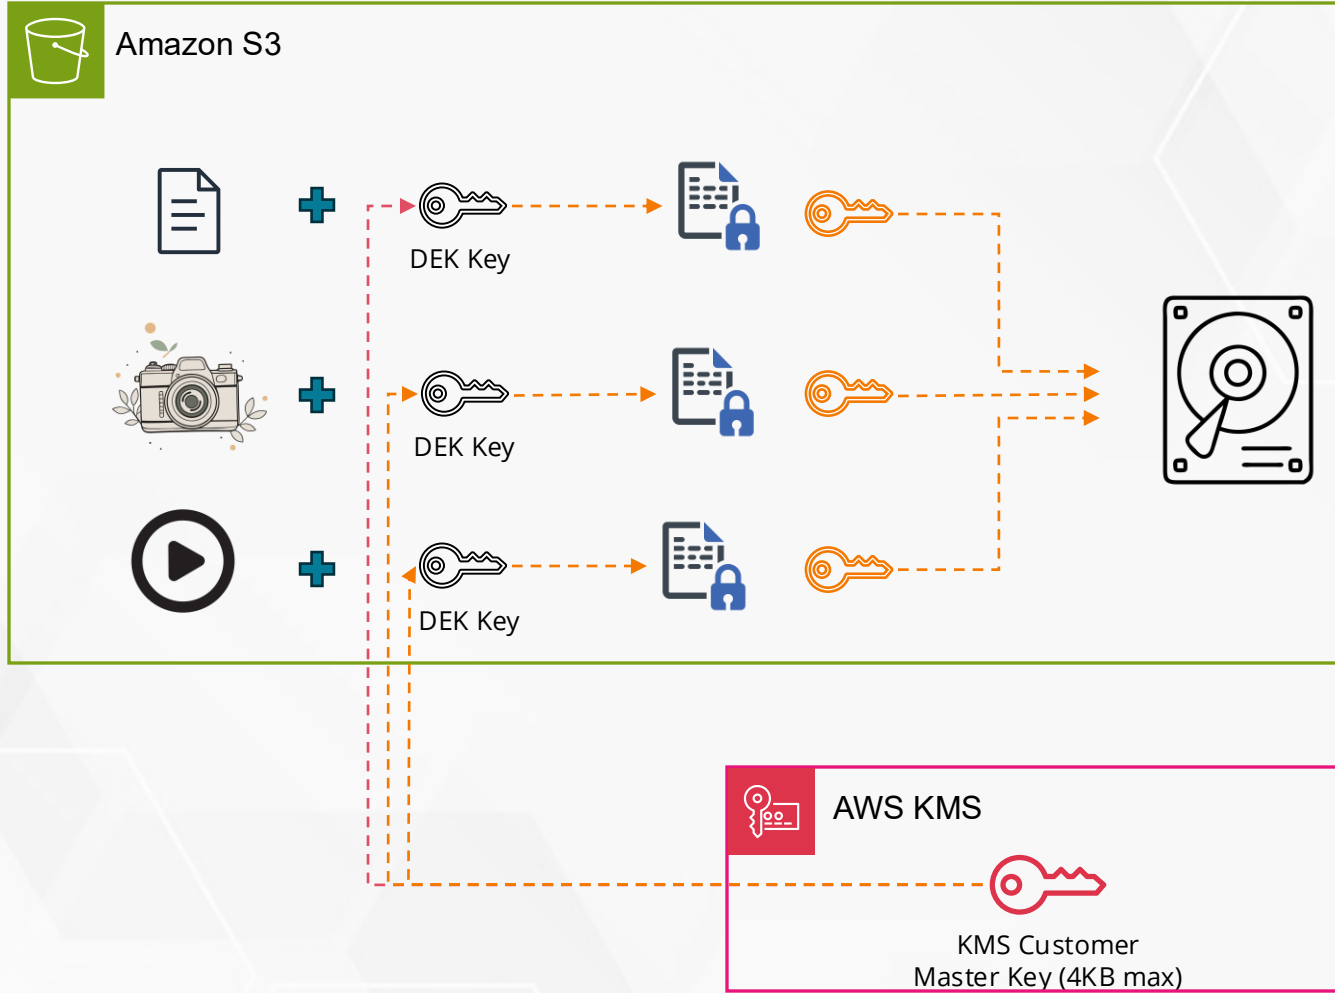

# Using Bucket Policies to Enforce Encryption

```
{
  "Version": "2012-10-17",
  "Id": "PutObjectPolicy",
  "Statement": [{
    "Sid": "DenyObjectsThatAreNotSSEKMS",
    "Effect": "Deny",
    "Principal": "*",
    "Action": "s3:PutObject",
    "Resource": "arn:aws:s3:::amzn-s3-demo-bucket1/*",
    "Condition": {
      "Null": {
        "s3:x-amz-server-side-encryption-aws-kms-key-id": "true"
      }
    }
  }]
}
```

You can specify the KMS key using the **x-amz-server-side-encryption-aws-kms-key-id** header or rely on your default bucket encryption configuration. If your **PutObject** request specifies **aws:kms** in the **x-amz-server-side-encryption** header but does not specify the **x-amz-server-side-encryption-aws-kms-key-id** header, then Amazon S3 assumes that you want to use the AWS managed key.

Uploads denied unless request includes the  
x-amz-server-side-encryption-aws-kms-key-id

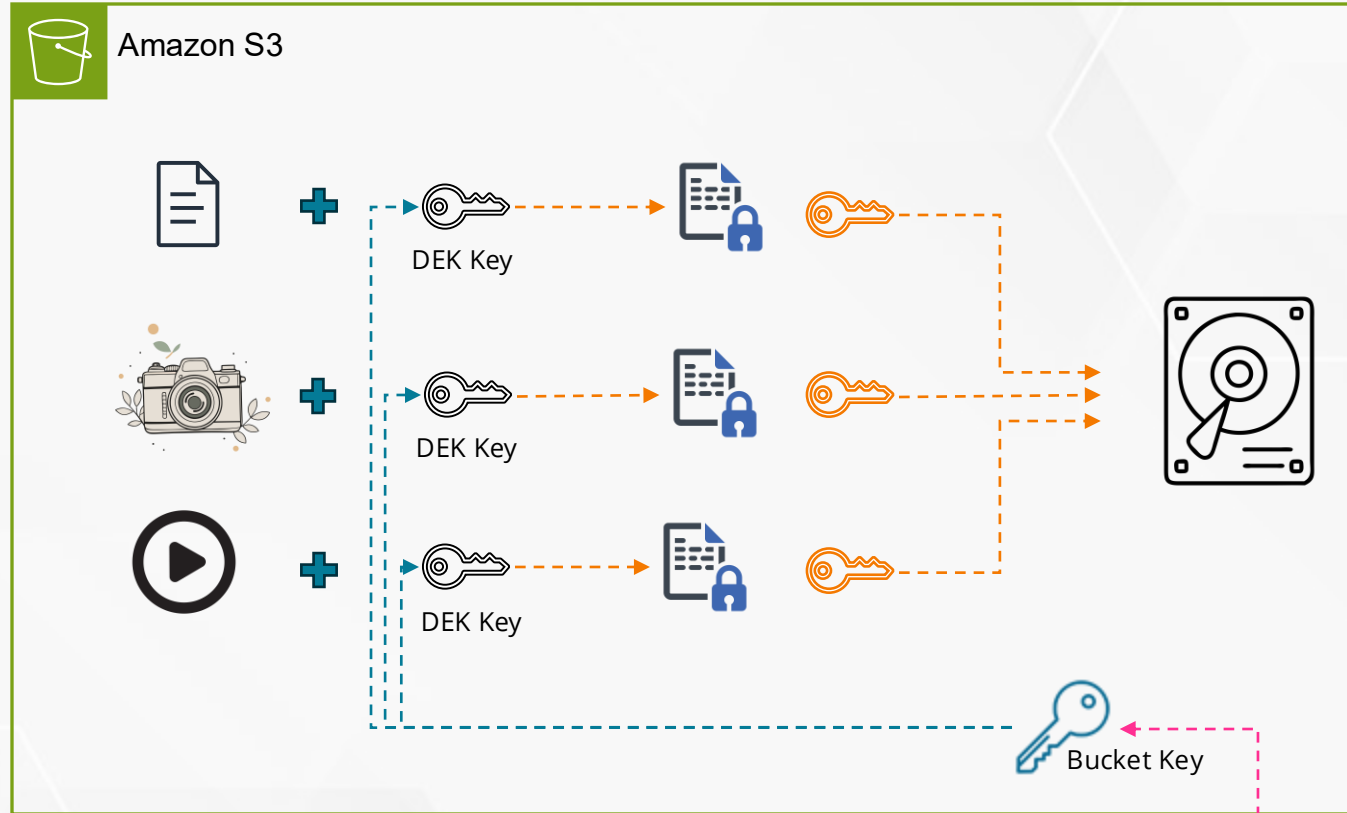
# Amazon S3 KMS Requests




## KMS Throttling

You have exceeded the rate at which you may call KMS. Reduce the frequency of your calls.  
(Service: AWSKMS; Status Code: 400; Error Code: ThrottlingException; Request ID: <ID>)

# Amazon S3 Bucket Keys



S3 Bucket Keys aren't supported for dual-layer server-side encryption with AWS Key Management Service (AWS KMS) keys (DSSE-KMS)





# Amazon S3 Encryption

Testing Encryption Options for Amazon S3



# Amazon S3 Presigned URLs

Sharing Objects Securely

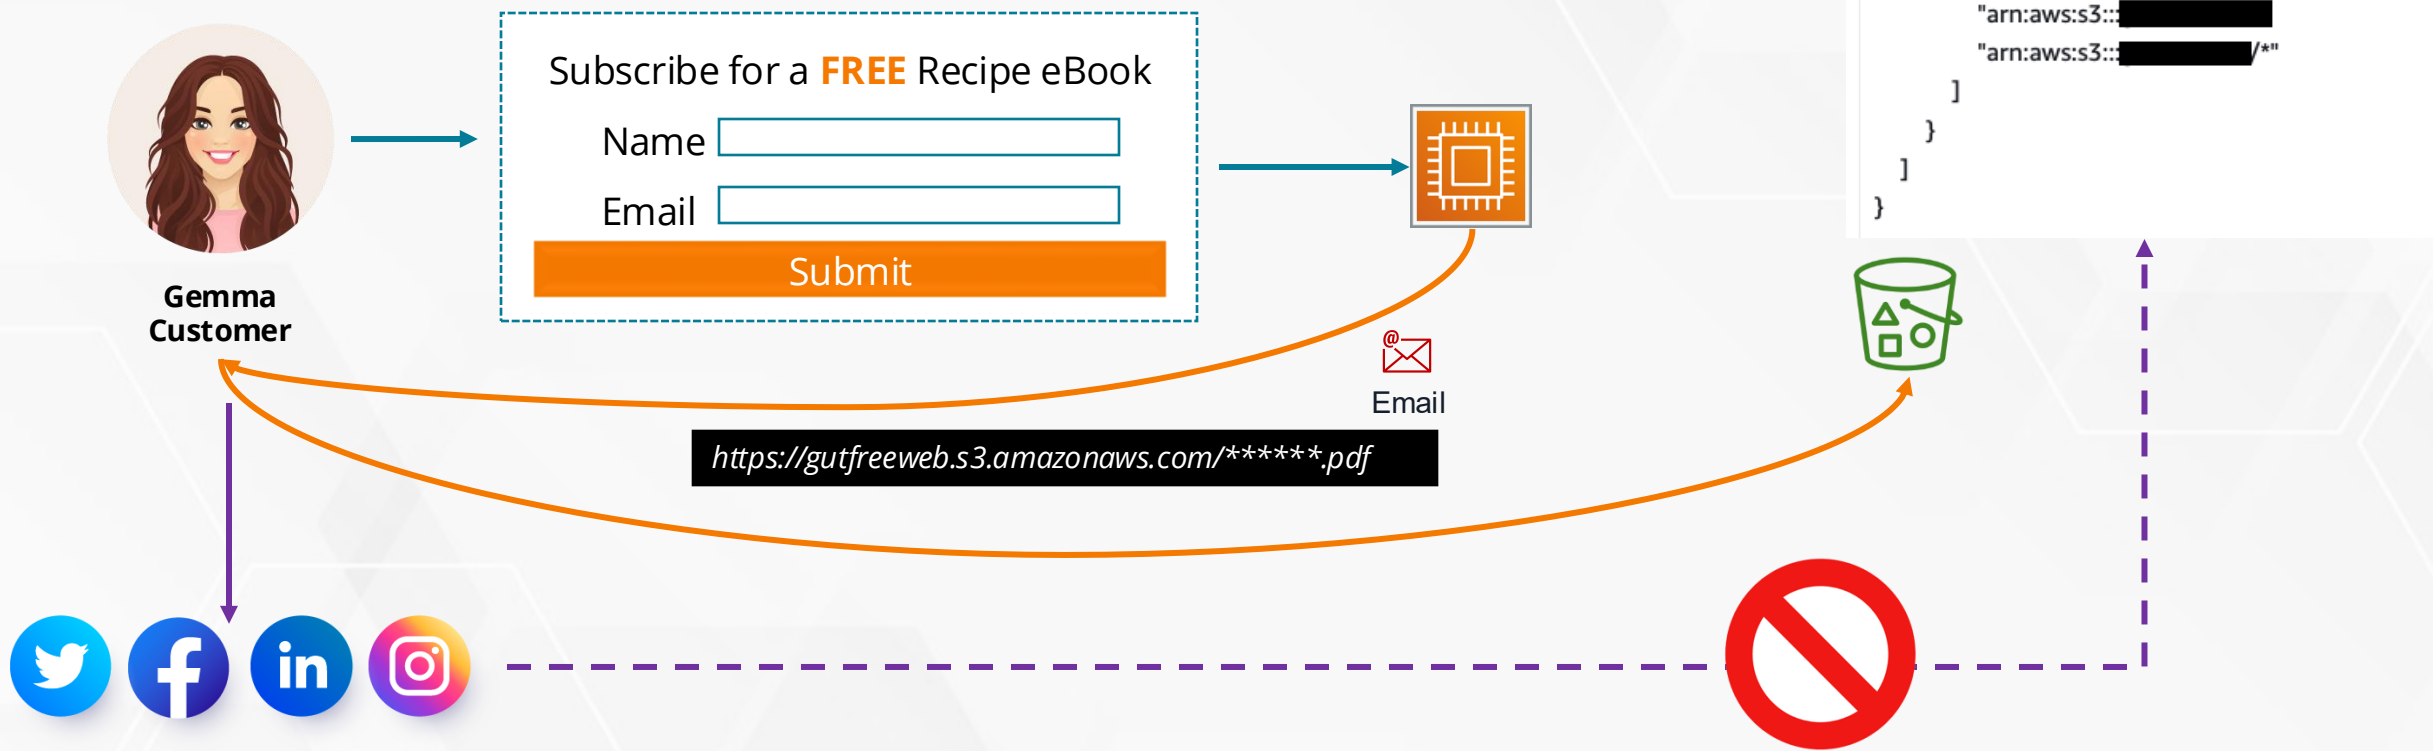


# Customer Scenario



# Fulfilling the business requirements

*“secure, scalable and resilient, ensure that the eBook is only available to genuine sign-ups.”*



```
{
  "Version": "2012-10-17",
  "Id": "Policy1705424029637",
  "Statement": [
    {
      "Sid": "Stmt1705424028548",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": [
        "arn:aws:s3:::[REDACTED]",
        "arn:aws:s3:::[REDACTED]/*"
      ]
    }
  ]
}
```

# Introducing Amazon S3 Pre-Signed URLs

*"secure, scalable and resilient, ensure that the eBook is only available to genuine sign-ups."*



Share "book.pdf" with a presigned URL

Presigned URLs are used to grant access to an object for a limited time. [Learn more](#)

Anyone can access the object with this presigned URL until it expires, even if the bucket, and object are private.

Time interval until the presigned URL expires

Using the S3 console, you can share an object with a presigned URL for up to 12 hours or until your session expires. To create a presigned URL with a longer time interval, use the AWS CLI or AWS SDK. Time intervals for presigned URLs can be restricted by your IAM policy.

☐ Minutes

☒ Hours

Number of hours

1

Must be a whole number between 1 and 12.

After you create the presigned URL, it's automatically copied to your clipboard.

Cancel Create presigned URL

**https://gutfree.s3.us-east-1.amazonaws.com/books/book.pdf?response-content-disposition=inline&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEGMaCWV1LXdlc3QtMijGMEQCIG5%2BalpeghxZHd%2BnjbmYQk%2FIdDN%%2FQN69gXQ6HvK8vE%2FRJzWCYOlh%2FxIN9WcAwlxjciT4jP05BgCpbUAh5RjOSMn%2FNKsOYsYX%2FWNhmidMsgIPTkK5AdTC1Gk5fdEl3H&X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Date=20240120T191219Z&X-Amz-SignedHeaders=host&X-Amz-Expires=3600&X-Amz-Credential=ASIA5XZWIW7TAEYOFEEEN%2F20240120%2Fus-east-1%2Fs3%2Faws4\_request&X-Amz-Signature=767af7db23ed4dbf5980d3e259bbebd61bd1b6dd4f7c0081a8ebdd0cff6e02e3**

# Presigned URLs

- When you create a presigned URL, you must provide your security credentials, and then specify the following:
- An Amazon S3 bucket
- An object key (if downloading this object will be in your Amazon S3 bucket, if uploading this is the file name to be uploaded)
- An HTTP method (GET for downloading objects or PUT for uploading)
- An expiration time interval



If you created a presigned URL using a temporary credential (such when using an IAM Role, the URL expires when the credential expires. In general, a presigned URL expires when the credential you used to create it is revoked, deleted, or deactivated. This is true even if the URL was created with a later expiration time



# Amazon S3 Server Access Logs

Logging Amazon S3 Activity

# Amazon S3 Server Access Logging





# Amazon S3 Object Lock

Fulfilling Compliance and Governance  
Requirements

# Overview of S3 Object Lock

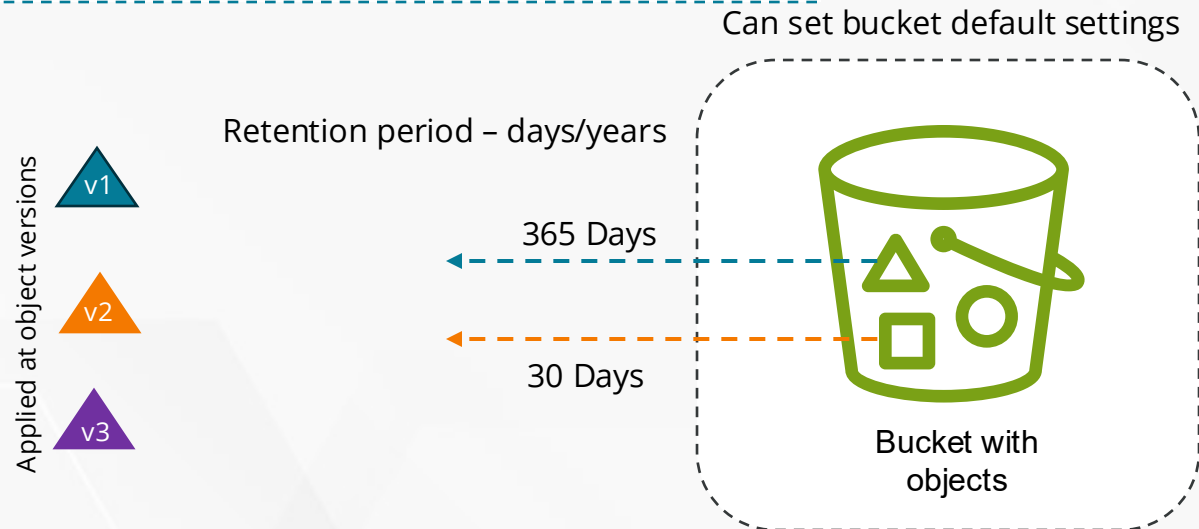
- S3 Object Lock can help prevent Amazon S3 objects from being deleted or overwritten for a fixed amount of time or indefinitely.
- Object Lock uses a write-once-read-many (WORM) mode
- Object Locks requires versioning to be enabled
- Two Options
  - S3 Object Lock **Retention period**
  - S3 Object Lock **Legal hold**

S3 stores lock information in metadata on the object version when locking that specific object version. The lock protects only the version that's specified in the request.



# Object Lock Retention period

Important Note: When you PUT an object version that has an explicit individual retention mode and period in a bucket, the object version's individual Object Lock settings override any bucket property retention settings



Restrict the minimum and maximum allowable retention periods with the s3:object-lock-remaining-retention-days condition key in the bucket policy

The only way to delete an object under the compliance mode before its retention date expires is to delete the associated AWS account.

## Compliance Mode



Object version can't be overwritten or deleted  
Retention mode can't be changed  
Retention period can't be shortened

```
aws s3api put-object-lock-configuration --bucket  
amzn-s3-demo-bucket1 --object-lock-configuration='{  
  "ObjectLockEnabled": "Enabled", "Rule": {  
    "DefaultRetention": { "Mode": "COMPLIANCE", "Days":  
      50 } } }
```



## Governance Mode




Cannot overwrite or delete an object version  
Cannot alter its lock settings  
**UNLESS** identity has special permissions

s3:BypassGovernanceRetention

Header: x-amz-bypass-governance-retention:true



# Object Lock – Legal holds



Place a *legal hold* on an object version

Legal holds are independent from retention periods

Object version is locked until legal hold is removed



placed and removed by any user  
who has s3:PutObjectLegalHold  
permissions



# Amazon S3 Select & Glacier Select


Filtering Data Retrievals

# Data Retrieval from Amazon S3

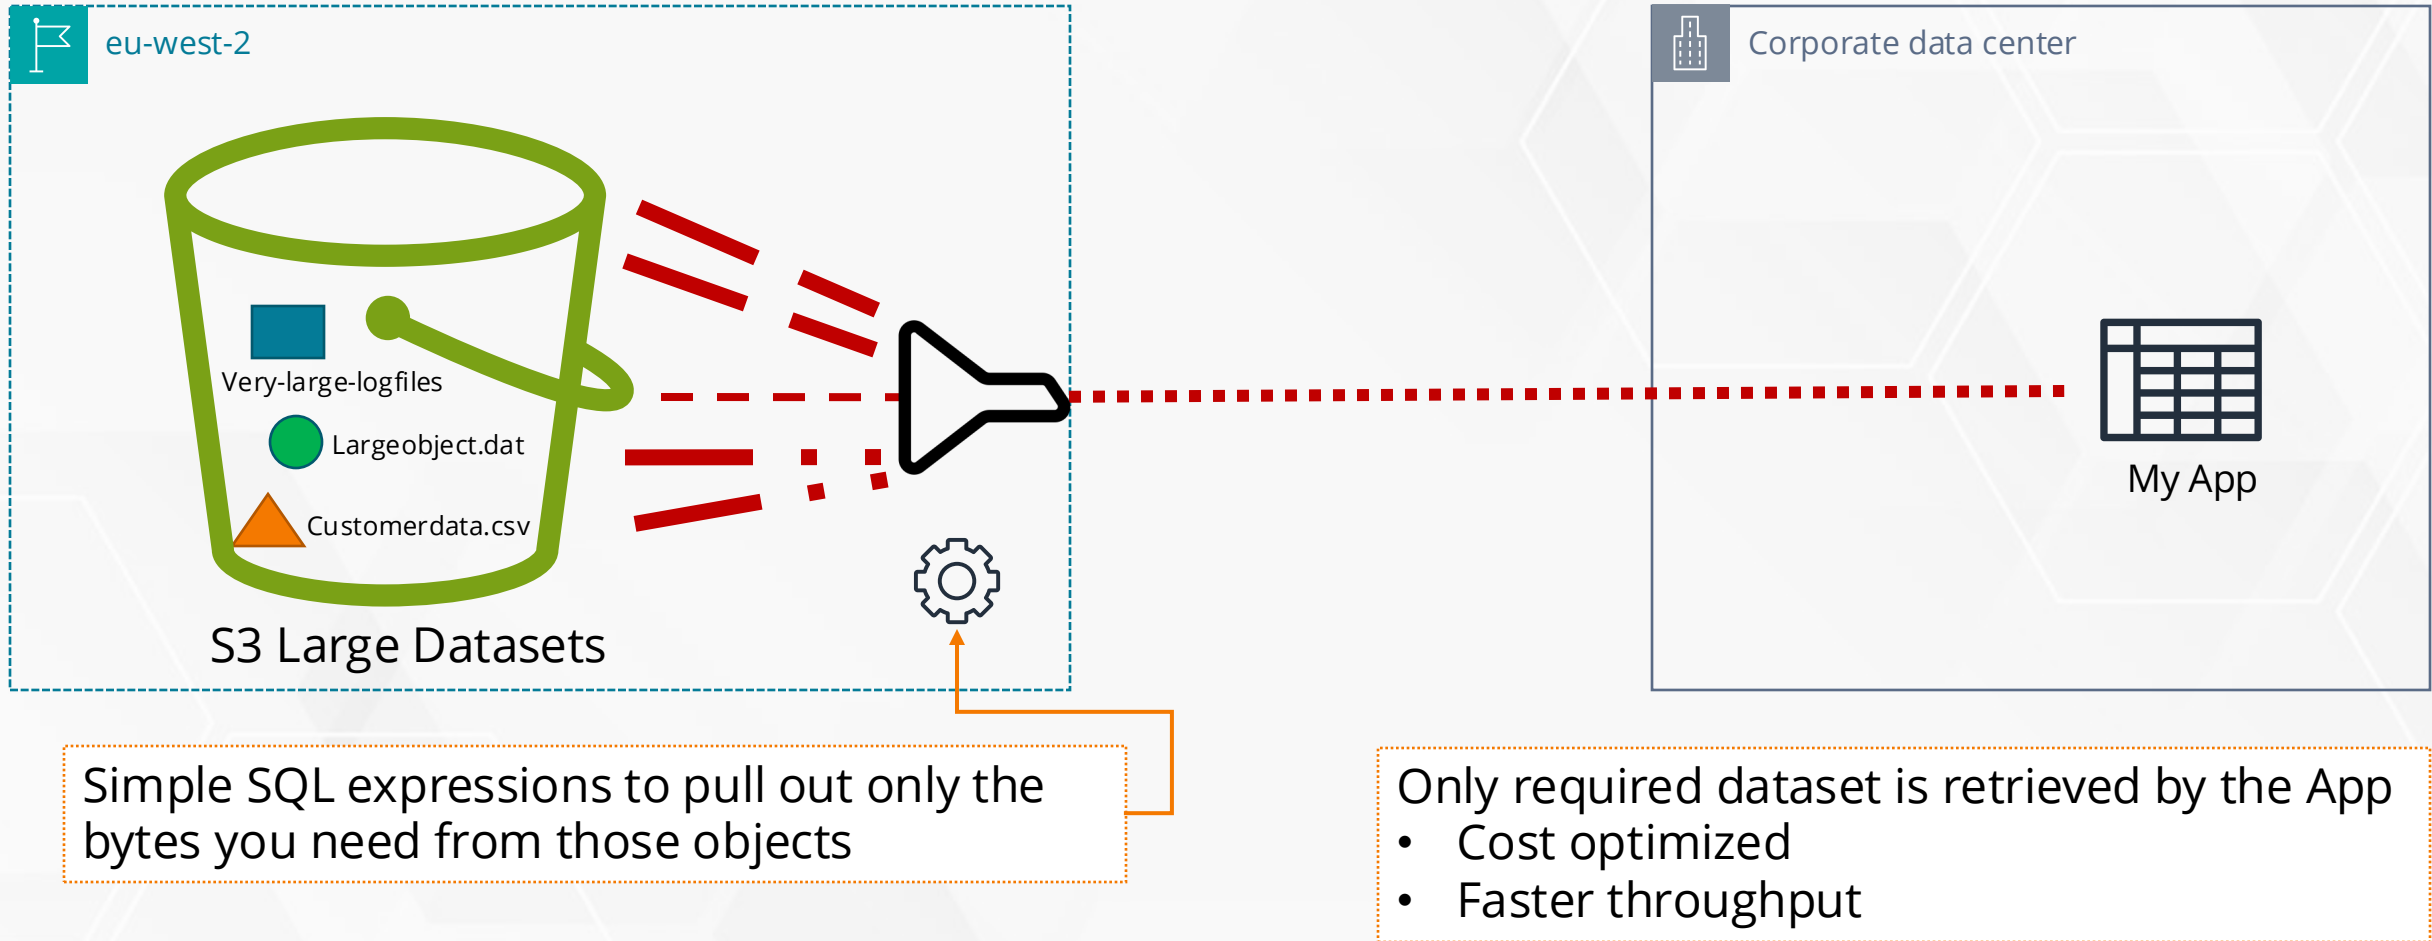
---

- Maximum Object Size is 5TB
- You typically retrieve objects in their entirety – e.g., 1GB CSV Customer Order across all stores
  - You consume 1GB of Data in data transfer
  - Time Delay in retrieval
  - Your application may filter data at client side, e.g., you only want to know the total breakdown of customers across five stores in Birmingham.
- S3 and Glacier Select allows you to filter data in the S3 environment before download using SQL-Like Statements, and only retrieve a subset of the required data

# Without Amazon S3 Select



# With Amazon S3 Select



# With Amazon S3 Glacier Select

